

Web Mining and Recommender Systems

Classification (& Regression Recap)

Learning Goals

In this section we want to:

- Explore techniques for **classification**
- Try some simple solutions, and see why they might fail
- Explore more complex solutions, and their advantages and disadvantages
- Understand the relationship between classification and regression
- Examine how we can reliably **evaluate** classifiers under different conditions


Recap...

Previously we started looking at
supervised learning problems

$$f(\text{data}) \xrightarrow{?} \text{labels}$$

Recap...

We studied **linear regression**, in order to learn linear relationships between features and parameters to predict **real-valued** outputs



Recap...

Pitch Black - Unrated Director's Cut R CC

★★★★★ 777 IMDb 7.1/10

[Watch Trailer](#)

When their ship crash-lands on a remote planet, the marooned passengers learn that escaped convict Riddick (Vin Diesel) isn't the only thing they have to fear. Deadly creatures lurk in the shadows, waiting to attack in the dark, and the planet is rapidly plunging into the

[See More](#)

Starring: Vin Diesel, Radha Mitchell
Runtime: 1 hour, 53 minutes
Available to watch on [supported devices](#).

Product Details

Genres	Science Fiction, Action, Horror
Director	David Twohy
Starring	Vin Diesel, Radha Mitchell
Supporting actors	Cole Hauser, Keith David, Lewis Fitz-Gerald, Claudia Black, Rhiana Granger, Angela Moore, Peter Chiang, Ken Twohy
Studio	NBC Universal
MPAA rating	R (Restricted)
Captions and subtitles	English Details
Rental rights	24 hour viewing period. Details
Purchase rights	Stream instantly and download to 2 locations Details
Format	Amazon Instant Video (streaming online video and digital download)

ratings
features

$f(\text{user features, movie features}) \xrightarrow{?} \text{star rating}$

Four important ideas:

1) Regression can be cast in terms of **maximizing a likelihood**

$$y_i = x_i \cdot \theta + \mathcal{N}(0, \sigma^2)$$

↓


$$d_i = x_i \cdot \theta - y_i$$

$$p_{\theta}(y_i | x_i)$$

Four important ideas:

2) Gradient descent for model optimization

1. Initialize θ at random
2. While (not converged) do
$$\theta := \theta - \alpha f'(\theta)$$



Four important ideas:

3) Regularization & Occam's razor

Regularization is the process of penalizing model complexity during training


$$\arg \min_{\theta} = \frac{1}{N} \|y - X\theta\|_2^2 + \lambda \|\theta\|_2^2$$


How much should we trade-off accuracy versus complexity?

Four important ideas:

4) Regularization pipeline

1. Training set – select model parameters
2. Validation set – to choose amongst models (i.e., hyperparameters)
3. Test set – just for testing!



for λ in $\{10^{-5} \dots 10^5\}$:

fit $\arg\min_{\theta} \text{error}(\theta) + \lambda \text{complexity}(\theta)$

$\text{valid}[\lambda] = \text{error}_{\text{valid}}(\theta)$

best = $\arg\min \text{valid}[\lambda]$

Model selection

A **validation set** is constructed to “tune” the model’s parameters

- Training set: used to **optimize the model’s parameters**
- Test set: used to report how well we expect the model to perform on **unseen data**
- Validation set: used to **tune** any model parameters that are not directly optimized


Model selection

A few “theorems” about training, validation, and test sets

- The training error **increases** as lambda **increases**
- The validation and test error are at least as large as the training error (assuming infinitely large random partitions)
- The validation/test error will usually have a “sweet spot” between under- and over-fitting

Up next...

How can we predict **binary** or **categorical** variables?



Up next...



Will I **purchase**
this product?
(yes)

Shop for engagement rings on Google

Sponsored ⓘ

 French-Set Halo Diamond... \$1,990.00 Ritani	 18K White Gold Delicate... \$950.00 Brilliant Earth ★★★★★ (57)	 18K White Gold Fancy D... \$1,825.00 Brilliant Earth ★★★★★ (13)	 Chamise Diamond Eng... \$975.00 Brilliant Earth ★★★★★ (7)
 Vintage Cushion Halo... \$4,140.00	 Princess Cut Diamond Eng... \$1,906.82	 18K White Gold Hudson... \$975.00	 18K White Gold Harmon... \$1,675.00

Will I **click on**
this ad?
(no)

Up next...

What animal appears in this image?
(mandarin duck)



Up next...

What are the **categories** of the item
being described?

(book, fiction, philosophical fiction)

From [Booklist](#)

Houellebecq's deeply philosophical novel is about an alienated young man searching for happiness in the computer age. Bored with the world and too weary to try to adapt to the foibles of friends and coworkers, he retreats into himself, descending into depression while attempting to analyze the passions of the people around him. Houellebecq uses his nameless narrator as a vehicle for extended exploration into the meanings and manifestations of love and desire in human interactions. Ironically, as the narrator attempts to define love in increasingly abstract terms, he becomes less and less capable of experiencing that which he is so desperate to understand. Intelligent and well written, the short novel is a thought-provoking inspection of a generation's confusion about all things sexual. Houellebecq captures precisely the cynical disillusionment of disaffected youth. *Bonnie Johnston* --This text refers to an out of print or unavailable edition of this title.

Up next...

We'll attempt to build **classifiers** that make decisions according to rules of the form

$$y_i = \begin{cases} 1 & \text{if } X_i \cdot \theta > 0 \\ 0 & \text{otherwise} \end{cases}$$

Up later...

1. Naïve Bayes

Assumes an **independence** relationship between the features and the class label and “learns” a simple model by counting

2. Logistic regression


Adapts the **regression** approaches we saw last week to binary problems

~~3.~~ Support Vector Machines

Learns to classify items by finding a hyperplane that separates them

Up later...

Ranking results in order of how likely they are to be relevant



tea station

Web Maps Shopping Images News More Search tools

About 20,900,000 results (0.61 seconds)

Tea Station 加州茶棧
teastationusa.com/ ▼
12 Tea Station locations in California and Nevada making Tea Station the ... We'd like to take this moment to thank you all tea lovers for your continued support.
3.8 ★★★★★ 19 Google reviews · Write a review · Google+ page


7315 Clairemont Mesa Boulevard, San Diego, CA 92111
(858) 268-8198
Menu - About - Ten Ren Products - San Gabriel

Tea Station - Kearny Mesa - San Diego, CA | Yelp
www.yelp.com > Restaurants > Chinese ▼ Yelp ▼

Up later...

Evaluating classifiers

- False positives are nuisances but false negatives are disastrous (or vice versa)
 - Some classes are very rare
- When we only care about the “most confident” predictions



e.g. which of these bags contains a weapon?

Web Mining and Recommender Systems

Classification: Naïve Bayes

Naïve Bayes

We want to associate a probability with a label and its negation:

$$p(\textit{label} | \textit{data})$$




$$p(\neg \textit{label} | \textit{data})$$

(classify according to whichever probability is greater than 0.5)

Q: How far can we get just by counting?

Naïve Bayes

e.g. $p(\text{movie is "action"} \mid \text{schwarzenegger in cast})$



Just count!

#films with Arnold = 45

#**action** films with Arnold = 32

$p(\text{movie is "action"} \mid \text{schwarzenegger in cast}) = 32/45$

Naïve Bayes

What about:

$p(\text{movie is "action" |}$
schwarzenegger in cast **and**
release year = 2017 **and**
mpaa rating = PG **and**
budget < \$1000000
)

#(training) fims with Arnold, released in 2017, rated PG, with a
budget below \$1M = 0

#(training) action fims with Arnold, released in 2017, rated PG,
with a budget below \$1M = 0

Naïve Bayes

Q: If we've never seen this combination of features before, what can we conclude about their probability?

A: We need some **simplifying assumption** in order to associate a probability with this feature combination

Naïve Bayes

Naïve Bayes assumes that features are **conditionally independent** given the label

$$(feature_i \perp\!\!\!\perp feature_j | label)$$

Naïve Bayes

$(feature_i \perp\!\!\!\perp feature_j | label)$

$$a \perp b \Rightarrow p(a, b) = p(a) p(b)$$

$$a \perp b | c \Rightarrow p(a, b | c) = p(a | c) p(b | c)$$

$a =$ you're wearing shorts

$b =$ I'm wearing shorts

$c =$ it's winter

Conditional independence?

$$(a \perp\!\!\!\perp b | c)$$

(a is conditionally independent of b, given c)

“if you know **c**, then knowing
a provides no additional
information about **b**”

(I remembered my umbrella $\perp\!\!\!\perp$ the streets are wet | it's raining)

Naïve Bayes

$$(feature_i \perp\!\!\!\perp feature_j | label)$$



$$p(feature_i, feature_j | label)$$

=

$$p(feature_i | label)p(feature_j | label)$$

Naïve Bayes

posterior prior likelihood

↓ ↓ ↓

$$p(\text{label}|\text{features}) = \frac{p(y) p(x|y)}{p(x)}$$

$P(y|x)$ $P(x)$

evidence

The diagram illustrates the Naïve Bayes formula. The word 'posterior' is positioned above the expression $p(\text{label}|\text{features})$, with a green arrow pointing down to it. Below this expression, the handwritten red text $P(y|x)$ is visible. To the right, the word 'prior' is above the handwritten red $p(y)$, with a green arrow pointing down to it. Further right, the word 'likelihood' is above the handwritten red $p(x|y)$, with a green arrow pointing down to it. Below the denominator $p(x)$, the word 'evidence' is written, with a green arrow pointing up to it. The entire equation is written in red ink.

Naïve Bayes

posterior

prior

likelihood

evidence

$$p(\text{label}|\text{features}) = \frac{p(\text{label})p(\text{features}|\text{label})}{p(\text{features})}$$

The diagram shows the components of Bayes' Theorem. 'posterior' points to $p(\text{label}|\text{features})$. 'prior' points to $p(\text{label})$. 'likelihood' points to $p(\text{features}|\text{label})$. 'evidence' points to $p(\text{features})$.

due to our conditional independence assumption:

$$p(label|features) = \frac{p(label) \prod_i p(feature_i|label)}{p(features)}$$

Naïve Bayes

$$p(\text{label}|\text{features}) = \frac{p(\text{label}) \prod_i p(\text{feature}_i|\text{label})}{\underbrace{p(\text{features})}}$$

vs.

$$p(\neg \text{label}|\text{features}) = \frac{p(\neg \text{label}) \prod_i p(\text{feature}_i|\neg \text{label})}{p(\text{features})}$$

The denominator doesn't matter, because we really just care about

$$p(\text{label}|\text{features}) \quad \text{vs.} \quad p(\neg \text{label}|\text{features})$$

both of which have the same denominator

Naïve Bayes

The denominator doesn't matter, because we really just care about

$$p(\textit{label}|\textit{features}) \quad \textbf{vs.} \quad p(\neg \textit{label}|\textit{features})$$

both of which have the same denominator

Web Mining and Recommender Systems

Naïve Bayes – Worked Example

Example 1

Amazon editorial descriptions:

Amazon.com Review

For most children, summer vacation is something to look forward to. But not for our 13-year-old uncle, and cousin who detest him. The third book in J.K. Rowling's [Harry Potter series](#) catapults Dursleys' dreadful visitor Aunt Marge to inflate like a monstrous balloon and drift up to the ceiling (and from officials at Hogwarts School of Witchcraft and Wizardry who strictly forbid students to go out into the darkness with his heavy trunk and his owl Hedwig).

As it turns out, Harry isn't punished at all for his errant wizardry. Instead he is mysteriously rescued by a triple-decker, violently purple bus to spend the remaining weeks of summer in a friendly inn called the Leaky Cauldron. In his third year at Hogwarts explains why the officials let him off easily. It seems that Sirius Black is loose. Not only that, but he's after Harry Potter. But why? And why do the Dementors, the guard dogs of the Ministry of Magic, are unaffected? Once again, Rowling has created a mystery that will have children and adults clamoring for the next book. Fortunately, there are four more in the works. (Ages 9 and older) --Karin Snelson --This text refers to the paperback edition.

50k descriptions:

Linked from course webpage

Example 1

P(book is a children's book |
 "wizard" is mentioned in the description **and**
 "witch" is mentioned in the description)

Code available on course webpage

Example 1

Conditional independence assumption:

"if you know **a book is for children**, then knowing that **wizards are mentioned** provides no additional information about whether **witches are mentioned**"

~~obviously~~ ridiculous

Double-counting

Q: What would happen if we trained two regressors, and attempted to “naively” combine their parameters?

Double-counting

$$\text{length} = \mathcal{O}_0 + \mathcal{O}_1[\text{wizards}]$$

$$\text{length} = \mathcal{O}_0 + \mathcal{O}_2[\text{witches}]$$

$$\stackrel{?}{\Rightarrow} \mathcal{O}_0 + \mathcal{O}_1[\text{wizards}] + \mathcal{O}_2[\text{witches}]$$

Double-counting

A: Since both features encode essentially the same information, we'll end up **double-counting** their effect

Web Mining and Recommender Systems

Classification: Logistic Regression

Learning Goals

- Introduce the **logistic regression** classifier
- Show how to design classifiers by maximizing a likelihood function

Logistic regression

Logistic Regression also aims
to model

$$p(\textit{label}|\textit{data})$$

By training a classifier of the
form

$$y_i = \begin{cases} 1 & \text{if } X_i \cdot \theta > 0 \\ 0 & \text{otherwise} \end{cases}$$

Logistic regression

Previously: regression

$$y_i = X_i \cdot \theta$$

Now: logistic regression


$$y_i = \begin{cases} 1 & \text{if } X_i \cdot \theta > 0 \\ 0 & \text{otherwise} \end{cases}$$

Logistic regression

Q: How to convert a real-valued expression ($X_i \cdot \theta \in \mathbb{R}$)
Into a probability
($p_\theta(y_i|X_i) \in [0, 1]$)

Logistic regression

A: sigmoid function: $\sigma(t) = \frac{1}{1+e^{-t}}$




Logistic regression

A: sigmoid function: $\sigma(t) = \frac{1}{1+e^{-t}}$

$P_{\theta}(y_i=1|x_i)$

$$\frac{1}{1+e^{-X_i \cdot \theta}}$$



Logistic regression

Training:

$X_i \cdot \theta$ should be maximized
when y_i is positive and
minimized when y_i is
negative

$$\arg \max_{\theta} \prod_{y_i=1} p_{\theta}(y_i=1|x_i) \prod_{y_i=0} (1 - p_{\theta}(y_i=1|x_i))$$

$\underbrace{\hspace{10em}}_{\sigma(x_i \cdot \theta)} \qquad \underbrace{\hspace{10em}}_{1 - \sigma(x_i \cdot \theta)}$

Logistic regression

Training:

$X_i \cdot \theta$ should be maximized
when y_i is positive and
minimized when y_i is
negative

$$\arg \max_{\theta} \prod_i \delta(y_i = 1) p_{\theta}(y_i | X_i) + \delta(y_i = 0) (1 - p_{\theta}(y_i | X_i))$$

 $\delta(\text{arg}) = 1$ if the argument is true, = 0 otherwise

Logistic regression

How to optimize?

$$L_{\theta}(y|X) = \prod_{y_i=1} p_{\theta}(y_i|X_i) \prod_{y_i=0} (1 - p_{\theta}(y_i|X_i))$$

- Take logarithm
- **Subtract** regularizer
- Compute gradient
- Solve using gradient **ascent**

Logistic regression

$$L_{\theta}(y|X) = \prod_{y_i=1} p_{\theta}(y_i|X_i) \prod_{y_i=0} (1 - p_{\theta}(y_i|X_i))$$

$$\begin{aligned} \ell_{\theta}(y|X) &= \sum_{y_i=1} \log \frac{1}{1+e^{-x_i \cdot \theta}} + \sum_{y_i=0} \log \frac{e^{-x_i \cdot \theta}}{1+e^{-x_i \cdot \theta}} \\ &= \sum_{y_i} -\log(1+e^{-x_i \cdot \theta}) + \sum_{y_i=0} -x_i \cdot \theta - \lambda \sum_k \theta_k^2 \end{aligned}$$

Logistic regression

$$l_{\theta}(y|X) = \sum_i -\log(1 + e^{-X_i \cdot \theta}) + \sum_{y_i=0} -X_i \cdot \theta - \lambda \|\theta\|_2^2$$

$$\frac{\partial l}{\partial \theta_k} = \sum_i \frac{x_{ik} e^{-x_i \cdot \theta}}{1 + e^{-x_i \cdot \theta}} + \sum_{y_i=0} -x_{ik} - 2\lambda \theta_k$$

⏟

$$\sum_i x_{ik} (1 - \sigma(x_i \cdot \theta))$$

Logistic regression

Log-likelihood:

$$l_{\theta}(y|X) = \sum_i -\log(1 + e^{-X_i \cdot \theta}) + \sum_{y_i=0} -X_i \cdot \theta - \lambda \|\theta\|_2^2$$

Derivative:

$$\frac{\partial l}{\partial \theta_k} = \sum_i X_{ik} (1 - \sigma(X_i \cdot \theta)) + \sum_{y_i=0} -X_{ik} - 2\lambda \theta_k$$

References

Further reading:

- On Discriminative vs. Generative classifiers: A comparison of logistic regression and naïve Bayes (Ng & Jordan '01)
- Boyd-Fletcher-Goldfarb-Shanno algorithm (BFGS)

Web Mining and Recommender Systems

Classification: Support Vector Machines

NOT COVERED!

So far we've seen...


So far we've looked at **logistic regression**, which is a classification model of the form:

$$y_i = \begin{cases} 1 & \text{if } X_i \cdot \theta > 0 \\ 0 & \text{otherwise} \end{cases}$$

- In order to do so, we made certain **modeling assumptions**, but there are many different models that rely on different assumptions
 - Next we'll look at another such model


(Rough) Motivation: SVMs vs Logistic regression

Q: Where would a logistic regressor place the decision boundary for these features?



SVMs vs Logistic regression

Q: Where would a logistic regressor place the decision boundary for these features?




SVMs vs Logistic regression

- Logistic regressors don't optimize the number of "mistakes"
- No special attention is paid to the "difficult" instances – every instance influences the model
- But "easy" instances can affect the model (and in a bad way!)
- How can we develop a classifier that optimizes the number of mislabeled examples?


Support Vector Machines: Basic idea

A classifier can be defined by the hyperplane (line) $\theta \mathbf{x} - \alpha = 0$




Support Vector Machines: Basic idea

Observation: Not all classifiers are equally good



Support Vector Machines



- An SVM seeks the classifier (in this case a line) that is **furthest from the nearest points**
- This can be written in terms of a specific optimization problem:


$$\arg \min_{\theta, \alpha} \frac{1}{2} \|\theta\|_2^2$$

such that

$$\forall_i y_i (\theta \cdot X_i - \alpha) \geq 1$$


Support Vector Machines

But: is finding such a separating hyperplane even possible?




Support Vector Machines

Or: is it actually a good idea?



Support Vector Machines

Want the margin to be as wide as possible



While penalizing points on the wrong side of it

Support Vector Machines

Soft-margin formulation:

$$\arg \min_{\theta, \alpha, \xi > 0} \frac{1}{2} \|\theta\|_2^2 + C \sum_i \xi_i$$

such that

$$\forall_i y_i (\theta \cdot X_i - \alpha) \geq 1 - \xi_i$$


Summary of Support Vector Machines

- SVMs seek to find a hyperplane (in two dimensions, a line) that optimally separates two classes of points
- The “best” classifier is the one that classifies all points correctly, such that the nearest points are **as far as possible** from the boundary
- If not all points can be correctly classified, a penalty is incurred that is proportional to **how badly the points are misclassified** (i.e., their distance from this hyperplane)

Web Mining and Recommender Systems

Classification – Worked example

Judging a book by its cover



[0.723845, 0.153926, 0.757238, 0.983643, ...]

4096-dimensional image features

Images features are available for each book on

http://cseweb.ucsd.edu/classes/fa19/cse258-a/data/book_images_5000.json



<http://caffe.berkeleyvision.org/>

Judging a book by its cover

Example: train a classifier to
predict whether a book is a
children's book from its cover
art

(code available on course webpage)

Judging a book by its cover


- The number of errors we made was extremely low, yet our classifier doesn't seem to be very good – why?
(stay tuned!)

Web Mining and Recommender Systems

Classifiers: Summary

Previously...

How can we predict **binary** or **categorical** variables?



Previously...



Will I **purchase**
this product?
(yes)

Shop for engagement rings on Google Sponsored ⓘ


 French-Set Halo Diamond... \$1,990.00 Ritani	 18K White Gold Delicate... \$950.00 Brilliant Earth ★★★★★ (57)	 18K White Gold Fancy D... \$1,825.00 Brilliant Earth ★★★★★ (13)	 Chamise Diamond Eng... \$975.00 Brilliant Earth ★★★★★ (7)
 Vintage Cushion Halo... \$4,140.00	 Princess Cut Diamond Eng... \$1,906.82	 18K White Gold Hudson... \$975.00	 18K White Gold Harmon... \$1,675.00

Will I **click on**
this ad?
(no)

Previously...

- **Naïve Bayes**
 - Probabilistic model (fits $p(\text{label}|\text{data})$)
 - Makes a conditional independence assumption of the form $(\text{feature}_i \perp\!\!\!\perp \text{feature}_j | \text{label})$ allowing us to define the model by computing $p(\text{feature}_i | \text{label})$ for each feature
 - Simple to compute just by counting
- **Logistic Regression**
 - Fixes the “double counting” problem present in naïve Bayes
- **SVMs**
 - Non-probabilistic: optimizes the classification error rather than the likelihood

1) Naïve Bayes



$$\begin{array}{ccc} \text{posterior} & & \text{prior} \quad \text{likelihood} \\ \downarrow & & \downarrow \quad \downarrow \\ p(\textit{label}|\textit{features}) = & \frac{p(\textit{label})p(\textit{features}|\textit{label})}{p(\textit{features})} \\ & \uparrow \\ & \text{evidence} \end{array}$$

due to our conditional independence assumption:

$$p(\textit{label}|\textit{features}) = \frac{p(\textit{label}) \prod_i p(\textit{feature}_i|\textit{label})}{p(\textit{features})}$$


2) logistic regression

sigmoid function: $\sigma(t) = \frac{1}{1+e^{-t}}$




Logistic regression

Q: Where would a logistic regressor place the decision boundary for these features?



Logistic regression

Q: Where would a logistic regressor place the decision boundary for these features?




Logistic regression

- Logistic regressors don't optimize the number of "mistakes"
- No special attention is paid to the "difficult" instances – every instance influences the model
- But "easy" instances can affect the model (and in a bad way!)
- How can we develop a classifier that optimizes the number of mislabeled examples?

3) Support Vector Machines

Can we train a classifier that optimizes the **number of mistakes**, rather than maximizing a probability?

Want the margin to be as wide as possible



While penalizing points on the wrong side of it

Pros/cons

- **Naïve Bayes**

- ++ Easiest to implement, most efficient to “train”
- ++ If we have a process that generates feature that *are* independent given the label, it’s a very sensible idea
- Otherwise it suffers from a “double-counting” issue

- **Logistic Regression**

- ++ Fixes the “double counting” problem present in naïve Bayes
- More expensive to train

- **SVMs**

- ++ Non-probabilistic: optimizes the classification error rather than the likelihood
- More expensive to train

Summary

- **Naïve Bayes**
 - Probabilistic model (fits $p(\text{label}|\text{data})$)
 - Makes a conditional independence assumption of the form $(\text{feature}_i \perp\!\!\!\perp \text{feature}_j | \text{label})$ allowing us to define the model by computing $p(\text{feature}_i | \text{label})$ for each feature
 - Simple to compute just by counting
- **Logistic Regression**
 - Fixes the “double counting” problem present in naïve Bayes
- **SVMs**
 - Non-probabilistic: optimizes the classification error rather than the likelihood


Web Mining and Recommender Systems

Evaluating classifiers

Which of these classifiers is best?

a:
errors = 2

b:
errors = 5



Which of these classifiers is best?


The solution which minimizes the
#errors may not be the best one

Which of these classifiers is best?

1. When data are highly imbalanced

If there are far fewer positive examples than negative examples we may want to assign additional weight to negative instances (or vice versa)


e.g. will I purchase a product? If I purchase 0.00001% of products, then a classifier which just predicts "no" everywhere is 99.99999% accurate, but not very useful



Which of these classifiers is best?

2. When mistakes are more costly in one direction

False positives are nuisances but false negatives are disastrous (or vice versa)




e.g. which of these bags contains a weapon?

Which of these classifiers is best?

3. When we only care about the “most confident” predictions

e.g. does a relevant result appear among the first page of results?



tea station

Web Maps Shopping Images News More Search tools

About 20,900,000 results (0.61 seconds)

Tea Station 加州茶棧
teastationusa.com/ ▾
12 Tea Station locations in California and Nevada making Tea Station the ... We'd like to take this moment to thank you all tea lovers for your continued support.
3.8 ★★★★★ 19 Google reviews · Write a review · Google+ page


7315 Clairemont Mesa Boulevard, San Diego, CA 92111
(858) 268-8198
Menu · About · Ten Ren Products · San Gabriel

Tea Station - Kearny Mesa - San Diego, CA | Yelp
www.yelp.com › Restaurants › Chinese ▾ Yelp ▾
★★★★★ Rating: 3 - 678 reviews - Price range: \$
678 Reviews of Tea Station "Taro tea with boba was soooo good! Great service, too! The shaved ice is very good at a reasonable price too."


Tea Station - Mira Mesa - San Diego, CA | Yelp
www.yelp.com › Restaurants › Taiwanese ▾ Yelp ▾
★★★★★ Rating: 3 - 381 reviews - Price range: \$
381 Reviews of Tea Station "Yes, I agree with Messiah! Everything is expensive but honestly the teas and boba are really delicious! But expect to wait long, the ...

Tea Station - Artesia, CA | Yelp
www.yelp.com › Food › Desserts ▾ Yelp ▾
★★★★★ Rating: 3.5 - 494 reviews - Price range: \$
494 Reviews of Tea Station "Came here at 12am SUPER hungry after not eating dinner. I was afraid the kitchen was going to be closed since they close at 1 am."

Evaluating classifiers




Evaluating classifiers




TP (true positive): Labeled as T , predicted as T

Evaluating classifiers




TN (true negative): Labeled as F , predicted as F

Evaluating classifiers



FP (false positive): Labeled as F , predicted as T

Evaluating classifiers



FN (false negative): Labeled as T , predicted as F

Evaluating classifiers

		Label	
		true	false
Prediction	true	true positive	false positive
	false	false negative	true negative

Classification accuracy = correct predictions / #predictions
= $(TP + TN) / (TP + TN + FP + FN)$

Error rate = incorrect predictions / #predictions
= $(FP + FN) /$ "

Evaluating classifiers

		Label	
		true	false
Prediction	true	true positive	false positive
	false	false negative	true negative

True positive rate (**TPR**) = true positives / #labeled positive

$$= \frac{TP}{TP + FN}$$

True negative rate (**TNR**) = true negatives / #labeled negative

$$= \frac{TN}{TN + FP}$$

Evaluating classifiers

		Label	
		true	false
Prediction	true	true positive	false positive
	false	false negative	true negative

$$1 - \frac{1}{2}(TPR + TNR)$$

$\swarrow$

$$\text{Balanced Error Rate (BER)} = \frac{1}{2} (\text{FPR} + \text{FNR})$$

= $\frac{1}{2}$ for a random/naïve classifier, 0 for a perfect classifier

Evaluating classifiers

e.g.

$\mathbf{y} = [1, -1, 1, 1, 1, -1, 1, 1, -1, 1]$
Confidence = $[1.3, -0.2, -0.1, -0.4, 1.4, 0.1, 0.8, 0.6, -0.8, 1.0] \rightarrow x.0$
1 -1 -1 -1 1 1 1 1 -1 1
TP TN FN FN TP FP TP TP TN TP

$$\begin{aligned}\# TP &= 5 \\ TN &= 2\end{aligned}$$

$$\begin{aligned}FP &= 1 \\ FN &= 2\end{aligned}$$

$$\begin{aligned}TPR &= \frac{5}{5+2} \\ TNR &= \frac{2}{2+1}\end{aligned}$$

$$BER = 1 - \frac{1}{2} \left(\frac{5}{7} + \frac{2}{3} \right)$$

Evaluating classifiers

How to optimize a balanced error measure:

$$L_{\theta}(y|X) = \prod_{y_i=1} p_{\theta}(y_i|X_i) \prod_{y_i=0} (1 - p_{\theta}(y_i|X_i))$$


$$l_{\theta}(y|X) = \frac{1}{|y_i=1|} \sum_{y_i=1} \log \sigma(x_i; \theta) + \frac{1}{|y_i=0|} \sum_{y_i=0} \log (1 - \sigma(x_i; \theta))$$

BFER

→ # of +ve instances

Evaluating classifiers – ranking

The classifiers we've seen can associate **scores** with each prediction



Evaluating classifiers – ranking

The classifiers we've seen can associate **scores** with each prediction

- In ranking settings, the actual labels assigned to the points (i.e., which side of the decision boundary they lie on) **don't matter**
- All that matters is that positively labeled points tend to be at **higher ranks** than negative ones

Evaluating classifiers – ranking

The classifiers we've seen can associate **scores** with each prediction

- For naïve Bayes, the "score" is the ratio between an item having a positive or negative class
- For logistic regression, the "score" is just the probability associated with the label being 1
- For Support Vector Machines, the score is the distance of the item from the decision boundary (together with the sign indicating what side it's on)

$x_i \cdot \theta$
or $\sigma(x_i \cdot \theta)$

Evaluating classifiers – ranking

The classifiers we've seen can associate **scores** with each prediction

e.g.

$\mathbf{y} = [1, -1, 1, 1, 1, -1, 1, 1, -1, 1]$
Confidence = $[1.3, -0.2, -0.1, -0.4, 1.4, 0.1, 0.8, 0.6, -0.8, 1.0]$

Sort **both** according to confidence:

$conf = [1.4, 1.3, 1.0, 0.8, 0.6, 0.1, -0.1, -0.2, -0.4, -0.8]$
 $y = [1, 1, 1, 1, 1, -1, -1, -1, -1, -1]$

Evaluating classifiers – ranking

The classifiers we've seen can associate **scores** with each prediction

Labels sorted by confidence:

[1, 1, 1, 1, 1, -1, 1, -1, 1, -1]



Suppose we have a fixed budget (say, six) of items that we can return (e.g. we have space for six results in an interface)

- Total number of **relevant** items =
- Number of items we returned =
- Number of **relevant items** we returned =

7
6
5

Evaluating classifiers – ranking

The classifiers we've seen can associate **scores** with each prediction

$y_i = T$   $x_i \cdot Q > 0$

$$\text{precision} = \frac{|\{\text{relevant documents}\} \cap \{\text{retrieved documents}\}|}{|\{\text{retrieved documents}\}|}$$

"fraction of retrieved documents that are relevant"

$$\text{recall} = \frac{|\{\text{relevant documents}\} \cap \{\text{retrieved documents}\}|}{|\{\text{relevant documents}\}|}$$

"fraction of relevant documents that were retrieved"

Evaluating classifiers – ranking

The classifiers we've seen can associate **scores** with each prediction

$\text{precision@}k$ = precision when we have a budget of k retrieved documents

the k highest values of $x_i \cdot \theta$

e.g.

- Total number of **relevant** items = 7
- Number of items we returned = 6 *— $k=6$*
- Number of **relevant items** we returned = 5

$$\text{precision@}6 = \frac{5}{6}$$

Evaluating classifiers – ranking

The classifiers we've seen can associate **scores** with each prediction

$$F_1 = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

(harmonic mean of precision and recall)

- Score $\rightarrow x; \theta$
- predict_proba
 $P_{\theta}(y|x) = \sigma(x; \theta)$


$$F_{\beta} = (1 + \beta^2) \cdot \frac{\text{precision} \cdot \text{recall}}{\beta^2 \text{precision} + \text{recall}}$$

(weighted, in case precision is more important (low beta), or recall is more important (high beta))

Precision/recall curves

How does our classifier behave as we “increase the budget” of the number retrieved items?

- For budgets of size 1 to N , compute the precision and recall
- Plot the precision against the recall




Summary

1. When data are highly imbalanced

If there are far fewer positive examples than negative examples we may want to assign additional weight to negative instances (or vice versa)

e.g. will I purchase product? If I purchase 0.00001% of products, then a classifier which just predicts "no" everywhere is 99.99999% accurate, but not very useful

Compute the true positive rate and true negative rate, and the F₁ score




Summary

2. When mistakes are more costly in one direction

False positives are nuisances but false negatives are disastrous (or vice versa)

Compute “weighted” error measures that trade-off the precision and the recall, like the F_{β} score




e.g. which of these bags contains a weapon?

Summary

3. When we only care about the “most confident” predictions

e.g. does
result
among
page of results?

Compute the precision@k , and
plot the signature of precision
versus recall



Web Mining and Recommender Systems

Classifier Evaluation: Worked Example

Code example: bankruptcy data

We'll look at a simple dataset from the UCI repository:

<https://archive.ics.uci.edu/ml/datasets/Polish+companies+bankruptcy+data>

```
@relation '5year-weka.filters.unsupervised.instance.SubsetByExpression-Enot ismissing(ATT20)'
```

```
@attribute Attr1 numeric
```

```
@attribute Attr2 numeric
```

```
...
```

```
@attribute Attr63 numeric
```

```
@attribute Attr64 numeric
```

```
@attribute class {0,1}
```

```
@data
```

```
0.088238,0.55472,0.01134,1.0205,-
```

```
66.52,0.34204,0.10949,0.57752,1.0881,0.32036,0.10949,0.1976,0.096885,0.10949,1475.2,0.24742,1.8027,0.10949,0.077287,50.199,
```

```
1.1574,0.13523,0.062287,0.41949,0.32036,0.20912,1.0387,0.026093,6.1267,0.37788,0.077287,155.33,2.3498,0.24377,0.13523,1.449
```

```
3,571.37,0.32101,0.095457,0.12879,0.11189,0.095457,127.3,77.096,0.45289,0.66883,54.621,0.10746,0.075859,1.0193,0.55407,0.42
```

```
557,0.73717,0.73866,15182,0.080955,0.27543,0.91905,0.002024,7.2711,4.7343,142.76,2.5568,3.2597,0
```

Did the company go bankrupt?

Code on course webpage

Web Mining and Recommender Systems

Supervised Learning: Summary so far

So far: Regression

Product Details

Genres	Science Fiction, Action, Horror
Director	David Twohy
Starring	Vin Diesel, Radha Mitchell
Supporting actors	Cole Hauser, Keith David, Lewis Fitz-Gerald, Claudia Black, Rhiana Gr Angela Moore, Peter Chiang, Ken Twohy
Studio	NBC Universal
MPAA rating	R (Restricted)
Captions and subtitles	English Details ▾
Rental rights	24 hour viewing period. Details ▾
Purchase rights	Stream instantly and download to 2 locations Details ▾
Format	Amazon Instant Video (streaming online video and digital download)

A. Phillips
Reviewer ranking: #17,230,554

90% helpful
votes received on reviews
(151 of 167)

ABOUT ME
Enjoy the reviews...

ACTIVITIES
Reviews (16)
Public Wish List (2)
Listmania Lists (2)
Tagged Items (1)




HipCzech
Aficionado
Male, from Texas
Profile Page

Member Since:	Jul 12, 2014	HipCzech was last seen:
Points:	175	Today at 12:19 AM
Beers:	108	
Places:	6	
Posts:	smoother than all of	0
Likes Received:	0	
Trading:	0% 0	

How can we use **features** such as product properties and user demographics to make predictions about **real-valued** outcomes (e.g. star ratings)?

How can we prevent our models from **overfitting** by favouring simpler models over more complex ones?



How can we assess our decision to optimize a particular error measure, like the MSE?

So far: Classification


Next we adapted these ideas to **binary** or **multiclass** outputs



What animal is in this image?




Will I **purchase** this product?




Will I **click on** this ad?



Combining features using naïve Bayes models



Logistic regression



Support vector machines

So far: supervised learning

Given **labeled training data** of the form

$$\{(\text{data}_1, \text{label}_1), \dots, (\text{data}_n, \text{label}_n)\}$$

Infer the function

$$f(\text{data}) \overset{?}{\rightarrow} \text{labels}$$

So far: supervised learning

We've looked at two types of prediction algorithms:

Regression  $y_i = X_i \cdot \theta$

Classification  $y_i = \begin{cases} 1 & \text{if } X_i \cdot \theta > 0 \\ 0 & \text{otherwise} \end{cases}$

Further Reading

Further reading:

- “Cheat sheet” of performance evaluation measures:
<http://www.damienfrancois.be/blog/files/modelperfcheatsheet.pdf>
 - Andrew Zisserman’s SVM slides, focused on
computer vision:
<http://www.robots.ox.ac.uk/~az/lectures/ml/lect2.pdf>