

Assignment 1 tips...

Assignment 1 tips...

Assignment 1 tips...

Web Mining and Recommender Systems

Content and structure in recommendation

Learning Goals

- Study how to incorporate *features* into recommender systems
- Look at a few types of recommendation that don't conform to the basic settings we've studied so far
- Explore systems for online advertising

Content and Structure in Recommender Systems

So far, our basic design principle when designing recommenders was to uncover *interactions between users and items*

We essentially argued that *interactions* can replace *features* in recommendation scenarios

In what scenarios might this logic not hold?

Content and Structure in Recommender Systems

In what scenarios might this logic not hold?

- **Cold-start:** We may have only a limited amount of interaction data; our previous argument applied only in the limit if we have lots of interactions per user and per item
- **Interpretability:** Features may help to understand what the model is doing, even if they don't improve performance
- **Temporal evolution:** None of the previous arguments hold if preferences/properties are non-stationary!
- **Other settings:** Some settings just don't look like simple user/item interaction data like we've seen so far

Content and Structure in Recommender Systems

Basic plan:

- Introduce ***Factorization Machines***, a general-purpose tool for modeling interactions among features
- Study a few other recommendation settings that don't conform to current assumptions (online dating, bartering, social recommendation)
- Talk about algorithms for ad recommendation

Web Mining and Recommender Systems

Factorization Machines

Factorization Machines

Factorization Machines are a general-purpose technique to extend (latent-factor) recommender systems to include arbitrary features

Rather than modeling interactions between users and items, we can model interactions between any pair of features (user and item interactions are then just one of many)

Factorization Machines

Idea: Describe our interaction data in terms of a feature matrix, including the user ID, item ID, and other features associated with an interaction

$$x(u, i) = \left[\underbrace{6601006}_{\text{user } u} \dots \underbrace{0010000}_{\text{item } i} \dots 15 \dots \underbrace{001000}_{\text{day}} \right]$$

Factorization Machines

Idea: Describe our interaction data in terms of a feature matrix, including the user ID, item ID, and other features associated with an interaction

$$\begin{matrix} X = & \begin{bmatrix} 1000000 \dots 000100000 \dots 0001000 \dots 15.95 \\ 0001000 \dots 000000010 \dots 0001000 \dots 12.25 \\ 0100000 \dots 000100000 \dots 0000010 \dots 15.00 \\ 0000100 \dots 010000000 \dots 0010000 \dots 17.50 \\ 1000000 \dots 000000010 \dots 1000000 \dots 19.95 \\ \underbrace{0000100}_{\text{user}} \dots \underbrace{000010000}_{\text{item}} \dots \underbrace{0000010}_{\text{weekday}} \dots \underbrace{10.15}_{\text{price}} \end{bmatrix} & y = \begin{bmatrix} 5 \\ 4 \\ \vdots \\ 2 \end{bmatrix} \end{matrix} \quad (\text{etc.})$$

Factorization Machines

Recall: We argued before why we can't just put all of these features into a simple linear model – we need to capture *interactions*

$$y = \langle x_{\text{user}}, \theta_{\text{user}} \rangle + \langle x_{\text{item}}, \theta_{\text{item}} \rangle$$

The diagram illustrates the decomposition of a linear model into user and item interaction components. At the top, the equation $y = x \cdot \theta$ is written in red. A red arrow points down from this equation to a more complex expression. The expression is $y = \langle x_{\text{user}}, \theta_{\text{user}} \rangle + \langle x_{\text{item}}, \theta_{\text{item}} \rangle$. The first term, $\langle x_{\text{user}}, \theta_{\text{user}} \rangle$, is written in red and has a red arrow pointing to it from the top equation. The second term, $\langle x_{\text{item}}, \theta_{\text{item}} \rangle$, is also written in red and is enclosed in a red oval. The entire expression is written in red ink.

Factorization Machines

Recall: We argued before why we can't just put all of these features into a simple linear model – we need to capture *interactions*

Factorization Machine Equation

Idea: Every feature is associated with a low-dimensional vector:

$$f(u, i) = \underbrace{w_0 + \sum_j w_j x_j}_{x \cdot \theta} + \sum_j \sum_{k > j} x_j x_k \theta_j \cdot \theta_k$$

$x(u, i)_j$

Factorization Machine Equation

Idea: Every feature is associated with a low-dimensional vector:

$$f(x) = w_0 + \underbrace{\sum_{i=1}^F w_i x_i}_{\text{offset and bias terms}} + \underbrace{\sum_{i=1}^F \sum_{j=i+1}^F \langle \gamma_i, \gamma_j \rangle x_i x_j}_{\text{feature interactions}}$$

$$\sim x + \sum_i x_i \beta_i + \sum_i \sum_j x_i x_j \gamma_i \gamma_j$$

Factorization Machine Equation - example

$$X = \left[\underbrace{000}_{\text{user}} \mid \underbrace{000}_{u} \dots \underbrace{001}_{\text{item}} \mid \underbrace{0000}_{i} \right]$$

$$\alpha + \sum_j x_j \beta_j + \sum_j \sum_{k > j} x_j x_k \gamma_j \gamma_k$$

$$= \alpha + \beta_u + \beta_i + \gamma_u \cdot \gamma_i$$

Factorization Machine Equation - example

$$x = [\underbrace{0001}_{\text{user}} \underbrace{000 \dots 00100000}_{\text{item } i} \underbrace{1}_{\text{price } p(i)}]$$

$$\alpha + \sum_j x_j \beta_j + \sum_j \sum_{k > j} x_j x_k \gamma_j \gamma_k$$

$$= \underbrace{\alpha + \beta_u + \beta_i}_{\text{}} + p(i) \beta_{\text{price}} + \gamma_u \cdot \gamma_i + p(i) \gamma_u \gamma_{\text{price}} + p(i) \gamma_i \gamma_{\text{price}}$$

Factorization Machine Equation – weekday feature

$$\left[\underbrace{1000000 \dots 000000010}_{\text{user}} \dots \underbrace{000010000}_{\text{item}} \dots \underbrace{0000010}_{\text{weekday}} \dots \underbrace{19.95}_{\text{price}} \right]$$

$$\begin{aligned} \alpha + \beta_u + \beta_i + \beta_w + \sigma_u \sigma_i \\ + \sigma_u \cdot \sigma_w \\ + \sigma_i \cdot \sigma_w \end{aligned}$$

Factorization Machines: FISM

$$f(u, i) = \alpha + \beta_u + \beta_i + \frac{1}{N} \sum_{j \in I_u \setminus \{i\}} \delta_j \cdot \delta_j'$$

~~δ_u~~

Recall: FISM (from Recommender Systems slides)

Replace

$$f(u, i) = \alpha + \beta_u + \beta_i + \gamma_u \cdot \gamma_i$$

with

$$f(u, i) = \alpha + \beta_u + \beta_i + \frac{1}{|I_u \setminus \{i\}|} \sum_{j \in I_u \setminus \{i\}} \gamma'_j \cdot \gamma_i$$

Can be implemented with a Factorization Machine!

Factorization Machines: FISM

$$\chi = \left[\underbrace{0 \ 0 \ 0 \ 0 \ 1 \ 0 \ 0 \ 0 \ 0}_{\text{item}} \quad \underbrace{\frac{1}{N} \ 0 \ 0 \ \frac{1}{N} \ 0 \ 0 \ \frac{1}{N} \ 0 \ 0 \ \frac{1}{N}}_{\text{user history} \left(\times \frac{1}{N} \right)} \right]$$

i

$$\alpha + \beta_i + \frac{1}{n} \sum_{j \in I_n} \chi_i \chi_j$$

Factorization Machines

Factorization Machines generalize many of the techniques we've seen to build "latent factor"-style models

And many of the techniques we'll see later: as we explore sequential/temporal recommenders, many can also be implemented via factorization machines!

Web Mining and Recommender Systems

Factorization Machines in fastFM

Factorization Machines – fastFM

```
In [4]: data = list(parseData(dataDir + "goodreads_reviews_comics_graphic.json.gz"))
```

```
In [5]: random.shuffle(data)
```

```
In [6]: data[0]
```

```
Out[6]: {'book_id': '15799191',  
         'date_added': 'Sun Jun 30 06:12:24 -0700 2013',  
         'date_updated': 'Thu Jul 04 08:02:07 -0700 2013',  
         'n_comments': 0,  
         'n_votes': 0,  
         'rating': 4,  
         'read_at': 'Thu Jul 04 08:02:07 -0700 2013',  
         'review_id': '86b43a448463928ac5d7887b364e2fcd',  
         'started_at': 'Sun Jun 30 00:00:00 -0700 2013',  
         'user_id': '23852e2647c217deb24964aadb26be64',  
         'year': 2013}
```

```
In [7]: userIDs, itemIDs = {}, {}
```

```
for d in data:  
    u, i = d['user_id'], d['book_id']  
    if not u in userIDs: userIDs[u] = len(userIDs)  
    if not i in itemIDs: itemIDs[i] = len(itemIDs)  
  
nUsers, nItems = len(userIDs), len(itemIDs)
```

$0 \dots N_u - 1$

$0 \dots N_i - 1$

```
In [8]: nUsers, nItems
```

```
Out[8]: (59347, 89311)
```

Factorization Machines – fastFM

```
In [9]: X = scipy.sparse.lil_matrix((len(data), nUsers + nItems))
```

```
In [10]: for i in range(len(data)):
          user = userIDs[data[i]['user_id']]
          item = itemIDs[data[i]['book_id']]
          X[i,user] = 1 # One-hot encoding of user
          X[i,nUsers + item] = 1 # One-hot encoding of item
```

Target (rating) to predict for each row

```
In [11]: y = numpy.array([d['rating'] for d in data])
```

Initialize the factorization machine

```
In [12]: fm = als.FMRegression(n_iter=1000, init_stdev=0.1, rank=5, l2_reg_w=0.1, l2_reg_V=0.5)
```

Split data into train and test portions

```
In [13]: X_train,y_train = X[:400000],y[:400000]
          X_test,y_test = X[400000:],y[400000:]
```

$[6010000 \quad 0100; 6000]$

Factorization Machines – fastFM

Train the model

```
In [14]: fm.fit(X_train, y_train)
```

```
Out[14]: FMRegression(init_stdev=0.1, l2_reg=0, l2_reg_V=0.5, l2_reg_w=0.1, n_iter=1000,  
                    random_state=123, rank=5)
```

Extract predictions on the test set

```
In [15]: y_pred = fm.predict(X_test)
```

```
In [16]: y_pred[:10]
```

```
Out[16]: array([2.23866277, 4.37847526, 3.84866594, 5.03002627, 3.94993096,  
                4.36740166, 4.22497778, 4.30029268, 3.28282377, 4.05036905])
```

Web Mining and Recommender Systems

Multisided Recommendation

Multisided recommendation

"Recommendation" so far has consisted of identifying which items a user is likely to "like best" or be most likely to interact with

When is this not suitable?

Multisided recommendation

When is this not suitable?

- items are users (e.g. online dating)
- constraints (e.g. budgets / online advertising)
- social dynamics (e.g. group rec.)

1. Group Recommendation (Amer-Yahia et al. 2009)

Recommendations should be made to satisfy the preferences of multiple people in a group (e.g. recommending a movie to a group of people)

How can we determine whether an item is compatible with a group of users?

$$f(u, i) \rightsquigarrow f(g, i)$$

1. Group Recommendation (Amer-Yahia et al. 2009)

(a) the item should be, **on average**, compatible with the users:

(given a user-to-item compatibility function $f(u,i)$)

Average Compatibility:

$$f(g,i) = \frac{1}{N} \sum_{u \in g} f(u,i)$$

1. Group Recommendation (Amer-Yahia et al. 2009)

(a) the item should be, **on average**, compatible with the users:

(given a user-to-item compatibility function $f(u,i)$)

Average Compatibility: $rel(\mathcal{G}, i) = \frac{1}{|\mathcal{G}|} \sum_{u \in \mathcal{G}} f(u, i)$

1. Group Recommendation (Amer-Yahia et al. 2009)

(b) no user should **hate** the item (e.g. consider recommendation in cases where users have constraints)

Least misery:

$$f(g, i) = \min_u f(u, i)$$

1. Group Recommendation (Amer-Yahia et al. 2009)

(b) no user should **hate** the item (e.g. consider recommendation in cases where users have constraints)

Least misery: $rel(\mathcal{G}, i) = \min_{u \in \mathcal{G}} f(u, i)$

1. Group Recommendation (Amer-Yahia et al. 2009)

(c) users should have **consensus** about the quality of an item (regardless of the actual scores)

Average pairwise disagreement:

$$dis(\mathcal{G}, i) = \frac{2}{|\mathcal{G}|(|\mathcal{G}| - 1)} \sum_{(u,v) \in \mathcal{G}, u \neq v} |f(u, i) - f(v, i)|$$

Disagreement variance:

$$dis(\mathcal{G}, i) = \frac{1}{|\mathcal{G}|} \sum_{u \in \mathcal{G}} \left(f(u, i) - \underbrace{\frac{1}{|\mathcal{G}|} \sum_{v \in \mathcal{G}} f(v, i)}_{\text{average compatibility in } \mathcal{G}} \right)^2.$$

1. Group Recommendation (Amer-Yahia et al. 2009)

Amer-Yahia suggest that group recommendation should consist of a combination of the two elements (high **relevance** and low **disagreement**):

$$\mathcal{F}(\mathcal{G}, i) = w_1 \times rel(\mathcal{G}, i) + w_2 \times (1 - dis(\mathcal{G}, i))$$

(see more in paper re. how to actually select sets of items, and user studies re. what kinds of relevance/agreement functions people like)

2. Online Dating (Pizzato et al. 2010)


Unlike other recommendation settings, the "items" in online dating settings are other users

A typical goal might consist of e.g. estimating reciprocal compatibility:

$$\text{reciprocal compatibility}(u, v) = \frac{2}{f(u, v)^{-1} + f(v, u)^{-1}}$$

3. Bartering (Rappaz et al. 2017)

Bartering (i.e., item exchange) is similar:
recommendations are pairs of users, who must have
compatible items to trade



$$f(u, i, v, j) = \frac{1}{2} (f(u, v, i) + f(v, u, j))$$
$$f(u, v, i) = \overset{\leftarrow}{f(u, i)}_{\text{compat}} + \sum_{\text{social}} \overset{\leftarrow}{u, v}_{\text{compat}}$$

Web Mining and Recommender Systems

Socially-Aware Recommendation

Socially-aware recommendation

In a recommendation setting, we probably don't care about predicting the social network, but it can still be useful:

- Our friends are likely (?) to have similar preferences to us
- Knowing what our friends do could help to predict what we'll do next

Socially-aware recommendation

Roughly speaking, we now have two sources of data to "explain":

$$R = \begin{pmatrix} 5 & 3 & \dots & \cdot \\ 4 & 2 & & 1 \\ 3 & \cdot & & 3 \\ \cdot & 2 & & 4 \\ 1 & 5 & & \cdot \\ \vdots & & \ddots & \vdots \\ 1 & 2 & \dots & \cdot \end{pmatrix} \quad S = \begin{pmatrix} 1 & 0 & \dots & 1 \\ 0 & 0 & & 1 \\ \vdots & & \ddots & \vdots \\ 1 & 0 & \dots & 1 \end{pmatrix}$$

Socially-aware recommendation

Two general strategies to incorporate social links:

1. Social signals can act as a *regularizer*. That is, a user's latent representation should be somehow similar to that of their friends
2. Social signals help us to harvest more data: friends' actions live somewhere in between positive and negative feedback

Socially-aware recommendation

1. Social signals can act as a *regularizer*.
 - The user parameter γ_u should simultaneously be able to predict interactions and social links
 - In practice this means that if we have few *interactions* associated with a user, we can still fit their model (estimate γ_u) based on their friends

Socially-aware recommendation


Use a *joint objective*:

Socially-aware recommendation

Use a *joint objective*:

$$\underbrace{\sum_{(u,i) \in R} (r_{u,i} - \sigma(\gamma_u \cdot \gamma_i))^2}_{\text{rating prediction error}} + \lambda^{(\text{trust})} \underbrace{\sum_{(u,v) \in A} (a_{u,v} - \sigma(\gamma_u \cdot \gamma'_v))^2}_{\text{trust prediction error}} + \lambda \|\gamma\|_2^2$$

Socially-aware recommendation



Socially-aware recommendation

2. Social signals act as additional data from which to generate samples

Recall: When we built Bayesian Personalized Ranking we were interested in the relative strength of different signals, e.g. an interaction should rank higher than a non-interaction

Socially-aware recommendation

What do friends' interactions mean?

(a) If a friend interacts with an item (that we haven't yet), we are more likely to interact with that item in the future

$$\underbrace{x_{u,i}}_{\text{positive}} \geq \underbrace{x_{u,k}}_{\text{social}} ; \quad \underbrace{x_{u,k}}_{\text{social}} \geq \underbrace{x_{u,j}}_{\text{negative}}$$

Socially-aware recommendation

What do friends' interactions mean?

(b) If a friend interacts with an item (that we haven't yet), we are *less* likely to interact with it in the future (why)?

$$\underbrace{x_{u,i}}_{\text{positive}} \geq \underbrace{x_{u,k}}_{\text{social}} ; \quad \underbrace{x_{u,i}}_{\text{positive}} \geq \underbrace{x_{u,j}}_{\text{negative}}$$

Web Mining and Recommender Systems

Recommendation in a few other settings

Music Recommendation

What's different about music?

- Hard to make use of signals: "listen to end" and "skip" are difficult to interpret as "positive" and "negative" (Pampalk et al. 2005)
 - A large fraction of consumption is *repeat consumption* (Anderson et al. 2014)
- Difficult to extract useful features from songs (e.g Wang and Wang, 2014)

Location-Based Social Networks

What role does geography play?

- Subsequent activities are likely to be nearby (see sequential recommendation, coming up!)
- Location features could be hierarchical (see e.g. Zheng et al. 2009)
- Other familiar problems, e.g. location-based networks can have a social component (etc.)

Price dynamics (lots in textbook)

Trivially, price can be incorporated into a factorization machine; but dynamics could be quite complicated

- Can we incorporate price *constraints* e.g. for recommending travel itineraries (Ge et al., 2011)
- Predict eventual purchase price from browsing trajectories (Hu et al., 2018)
- How do changes in price affect purchasing decisions (e.g. purchase quantity or brand) (Wan et al., 2017)

Temporal, Textual, Visual Features

Lots still to cover (also in textbook)


- Incorporating **text** into recommender systems is not so straightforward; note that (e.g.) review data isn't available at test time! Lots of other text-based tasks too
- **Visual** features: high-dimensional, relationship to user preferences is subtle (coming up later)
- **Temporal dynamics**: turn out to be critical in real-world recommender systems (coming up later)

Web Mining and Recommender Systems

Algorithms for advertising

Classification


Predicting which ads people click on might be a **classification** problem



Will I **click on**
this ad?

Recommendation

Or... predicting which ads people click on might be a **recommendation** problem



Advertising

So, we already have good algorithms for **predicting** whether a person would click on an ad, and generally for **recommending** items that people will enjoy.

So what's different about **ad recommendation**?

Advertising

1. We can't recommend everybody the same thing (even if they all want it!)

- Advertisers have a limited budget – they wouldn't be able to afford having their content recommended to everyone
- Advertisers **place bids** – we must take their bid into account (as well as the user's preferences – or not)
- In other words, we need to consider **both** what the **user and the advertiser** want (this is in contrast to recommender systems, where the content didn't get a say about whether it was recommended!)

2. We need to be **timely**

- We want to make a personalized recommendations immediately (e.g. the moment a user clicks on an ad) – this means that we can't train complicated algorithms (like what we saw with recommender systems) in order to make recommendations later
- We also want to update users' models **immediately** in response to their actions
 - (Also true for some recommender systems)


3. We need to take **context** into account

- Is the page a user is currently visiting particularly relevant to a particular type of content?
 - Even if we have a good model of the user, recommending them the same type of thing over and over again is unlikely to succeed – nor does it teach us anything **new** about the user
- In other words, there's an **explore-exploit** tradeoff – we want to recommend things a user will enjoy (exploit), but also to discover new interests that the user may have (explore)

Advertising

So, ultimately we need


- 1) Algorithms to match users and ads, given **budget constraints**



Advertising

So, ultimately we need

2) Algorithms that work in real-time and don't depend on monolithic optimization problems



Advertising

So, ultimately we need

- 3) Algorithms that adapt to users and capture the notion of an exploit/explore tradeoff



Web Mining and Recommender Systems

Advertising: Matching problems

Let's start with...

1. We can't recommend everybody the same thing (even if they all want it!)

- Advertisers have a limited budget – they wouldn't be able to afford having their content recommended to everyone
- Advertisers **place bids** – we must take their bid into account (as well as the user's preferences – or not)
- In other words, we need to consider **both** what the **user and the advertiser** want (this is in contrast to recommender systems, where the content didn't get a say about whether it was recommended!)

Bipartite matching

Let's start with a simple version of the problem we ultimately want to solve:

- 1) Every advertiser wants to show **one ad**
- 2) Every user gets to see **one ad**
- 3) We have some pre-existing model that assigns a score to user-item pairs

Bipartite matching


Suppose we're given some scoring function:

$$f(u, a) = \text{score for showing user } u \text{ ad } a$$

Could be:

- How much the owner of **a** is willing to pay to show their ad to **u**
- How much we expect the user **u** to spend if they click the ad **a**
- Probability that user **u** will click the ad **a**

Output of a regressor / logistic regressor!




Bipartite matching

Then, we'd like to show each user one ad, and we'd like each ad to be shown exactly once **so as to maximize this score** (bids, expected profit, probability of clicking etc.)

$$\sum_u f(u, ad(u))$$

s.t.

$$ad(u) = ad(v) \rightarrow u = v$$



each advertiser gets to show one ad

Bipartite matching

Then, we'd like to show each user one ad, and we'd like each ad to be shown exactly once **so as to maximize this score** (bids, expected profit, probability of clicking etc.)

$$\sum_{u,a} A_{u,a} f(u, a)$$

s.t.


$$\forall a \sum_u A_{u,a} = 1$$

 each advertiser gets to show one ad

Bipartite matching

We can set this up as a **bipartite matching** problem


- Construct a complete bipartite graph between users and ads, where each edge is weighted according to $f(u, a)$
- Choose edges such that each node is connected to exactly one edge



Bipartite matching

This is similar to the problem solved by (e.g.) online dating sites
to match men to women

For this reason it is called a **marriage problem**



Bipartite matching

This is similar to the problem solved by (e.g.) online dating sites
to match men to women

For this reason it is called a **marriage problem**

- A group of men should marry an (equally sized) group of women such that happiness is maximized, where “happiness” is measured by $f(m,w)$



compatibility between male m and female w


- Marriages are monogamous, heterosexual, and everyone gets married

(see also the original formulation, in which men have a preference function over women, and women have a *different* preference function over men)

Bipartite matching

We'll see one solution to this problem,
known as **stable marriage**

- Maximizing happiness turns out to be quite hard
 - **But**, a solution is "**unstable**" if:



- A man m is matched to a woman w' but would prefer w (i.e., $f(m, w') < f(m, w)$)


and

- The feeling is mutual – w prefers m to her partner (i.e., $f(w, m') < f(m, w)$)
- In other words, m and w would both want to "cheat" with each other

Bipartite matching

We'll see one solution to this problem,
known as **stable marriage**

- A solution is said to be **stable** if this is **never satisfied** for any pair (m, w)



- Some people may covet another partner,
but
- The feeling is never reciprocated by the other person
- So no pair of people would **mutually** want to cheat

Bipartite matching

The algorithm works as follows:

(due to Lloyd Shapley & Alvin Roth)

- Men propose to women (this algorithm is from 1962!)
- While there is a man m who is not engaged
 - He selects his most compatible partner, $\max_w f(m, w)$ (to whom he has not already proposed)
 - If she is not engaged, they become engaged
 - If she *is* engaged (to m'), but prefers m , she breaks things off with m' and becomes engaged to m instead

Bipartite matching

The algorithm works as follows:
(due to Lloyd Shapley & Alvin Roth)

```
All men and all women are initially 'free' (i.e., not engaged)
while there is a free man m, and a woman he has not proposed to
    w = max_w f(m,w)
    if (w is free):
        (m,w) become engaged (and are no longer free)
    else (w is engaged to m'):
        if w prefers m to m' (i.e.,  $f(m,w) > f(m',w)$ ):
            (m,w) become engaged
            m' becomes free
```

Bipartite matching

The algorithm works as follows:

(due to Lloyd Shapley & Alvin Roth)

- The algorithm terminates

Bipartite matching

The algorithm works as follows:

(due to Lloyd Shapley & Alvin Roth)

- The algorithm terminates

(either the number of free people decreases at each step, or, if it stays the same, the happiness increases)

Bipartite matching

The algorithm works as follows:

(due to Lloyd Shapley & Alvin Roth)

- The solution is stable

Bipartite matching

The algorithm works as follows:

(due to Lloyd Shapley & Alvin Roth)

- The solution is stable

(suppose m and w prefer each other to their current partners, w' and m')

But m would have proposed to w before he proposed to w'

- if w rejected his proposal, she must have been with someone she liked better
- if w accepted his proposal (but dumped him later), it must also have been for someone she likes better)

Bipartite matching

The algorithm works as follows:

(due to Lloyd Shapley & Alvin Roth)

- The solution is $O(n^2)$

Bipartite matching

The algorithm works as follows:

(due to Lloyd Shapley & Alvin Roth)

- The solution is $O(n^2)$

(every proposal is made at most once, and there are $O(n^2)$ proposals)

The input is $O(n^2)$ (i.e., the compatibility function) so it certainly couldn't be **better** than $O(n^2)$)

Bipartite matching – extensions/improvements

Can all of this be improved upon?

1) It's not optimal

Bipartite matching – extensions/improvements

Can all of this be improved upon?

1) It's not optimal


- Although there's no **pair** of individuals who would be happier by cheating, there could be **groups** of men and women who would be ultimately happier if the graph were rewired
- To get a **truly optimal** solution, there's a more complicated algorithm, known as the "Hungarian Algorithm"
 - But it's $O(n^3)$
- And really complicated and unintuitive (but there's a ref later)

Bipartite matching – extensions/improvements

Can all of this be improved upon?

2) Marriages are **monogamous**, heterosexual, and everyone gets married

(each user gets shown two ads, each ad gets shown to two users)




- Each advertiser may have a fixed budget of (1 or more) ads
- We may have room to show more than one ad to each customer
- See “Stable marriage with multiple partners: efficient search for an optimal solution” (refs)

Bipartite matching – extensions/improvements

Can all of this be improved upon?

2) Marriages are monogamous,
heterosexual, and everyone gets married




- This version of the problem is known as **graph cover** (select edges such that each node is connected to exactly one edge)
- The algorithm we saw is really just graph cover for a bipartite graph
- Can be solved via the “stable roommates” algorithm (see refs) and extended in the same ways

Bipartite matching – extensions/improvements

Can all of this be improved upon?

2) Marriages are monogamous,
heterosexual, and everyone gets married




- This version of the problem can address a very different variety of applications compared to the bipartite version
 - Roommate matching
 - Finding chat partners
- (or any sort of person-to-person matching)

Bipartite matching – extensions/improvements

Can all of this be improved upon?

2) Marriages are monogamous,
heterosexual, and **everyone gets married**



- Easy enough just to create “dummy nodes” that represent no match

no ad is shown to the corresponding user

Bipartite matching – applications


Why are matching problems so important?

- Advertising
 - Recommendation
 - Roommate assignments
 - Assigning students to classes
- General resource allocation problems
- Transportation problems (see “Methods of Finding the Minimal Kilometrage in Cargo-transportation in space”)
 - Hospitals/residents

Bipartite matching – applications

Why are matching problems so important?

- Point pattern matching



Bipartite matching – extensions/improvements

What about more complicated rules?

- (e.g. for hospital residencies) Suppose we want to keep couples together
- Then we would need a more complicated function that encodes these pairwise relationships:

$$\sum_{u,v} f(u, v, \text{hospital}(u), \text{hospital}(v))$$

pair of residents

hospitals to which they're assigned

So far...

Surfacing ads to users is a little like building a **recommender system** for ads

- We need to model the compatibility between each user and each ad (probability of clicking, expected return, etc.)
- **But**, we can't recommend the same ad to every user, so we have to handle "budgets" (both how many ads can be shown to each user and how many impressions the advertiser can afford)
- **So**, we can cast the problem as one of "covering" a bipartite graph
- Such **bipartite matching** formulations can be adapted to a wide variety of tasks

Questions?

Further reading:

- The original stable marriage paper

"College Admissions and the Stability of Marriage" (Gale, D.; Shapley, L. S., 1962):

<https://www.jstor.org/stable/2312726>

- The Hungarian algorithm

"The Hungarian Method for the assignment problem" (Kuhn, 1955):

<https://tom.host.cs.st-andrews.ac.uk/CS3052-CC/Practicals/Kuhn.pdf>

- Multiple partners

"Stable marriage with multiple partners: efficient search for an optimal solution" (Bansal et al., 2003)

- Graph cover & stable roommates

"An efficient algorithm for the 'stable roommates' problem" (Irving, 1985)

<https://dx.doi.org/10.1016%2F0196-6774%2885%2990033-1>

Web Mining and Recommender Systems

AdWords

Advertising

1. We can't recommend everybody the same thing (even if they all want it!)

- So far, we have an algorithm that takes "budgets" into account, so that users are shown a limited number of ads, and ads are shown to a limited number of users
- **But**, all of this only applies if we see all the users and all the ads **in advance**
 - This is what's called an **offline algorithm**


2. We need to be **timely**

- But in many settings, users/queries come in one at a time, and need to be shown some (highly compatible) ads
 - But we still want to satisfy the same quality and budget constraints
- So, we need **online algorithms** for ad recommendation

What is adwords?

Adwords allows advertisers to bid on keywords

- This is similar to our matching setting in that advertisers have limited **budgets**, and we have limited space to show ads




What is adwords?

Adwords allows advertisers to bid on keywords

- This is similar to our matching setting in that advertisers have limited **budgets**, and we have limited space to show ads
 - **But**, it has a number of key differences:
 1. Advertisers don't pay for impressions, but rather they pay when their ads get clicked on
 2. We don't get to see all of the queries (keywords) in advance – **they come one-at-a-time**

What is adwords?

Adwords allows advertisers to bid on keywords



- We still want to match advertisers to keywords to satisfy budget constraints
- But can't treat it as a monolithic optimization problem like we did before
- Rather, we need an **online** algorithm

What is adwords?

Suppose we're given

- Bids that each advertiser is willing to make for each query

$$\begin{array}{ccc} & f(q, a) & \\ \nearrow & & \nwarrow \\ \text{query} & & \text{advertiser} \end{array}$$

(this is how much they'll pay **if the ad is clicked on**)

- Each is associated with a click-through rate

$$ctr(q, a)$$

- Budget for each advertiser $b(a)$ (say for a 1-week period)
- A limit on how many ads can be returned for each query

What is adwords?

And, every time we see a query

- Return at most the number of ads that can fit on a page
 - And which won't overrun the budget of the advertiser (if the ad is clicked on)

Ultimately, what we want is an algorithm that maximizes **revenue** – the number of ads that are clicked on, multiplied by the bids on those ads

Competitiveness ratio

What we'd like is:

the revenue should be as close as possible to what we *would* have obtained if we'd seen the whole problem up front

(i.e., if we didn't have to solve it online)

We'll define the **competitive ratio** as:

$$\frac{\text{revenue of our algorithm}}{\text{revenue of an optimal algorithm}}$$

Greedy solution

Let's start with a simple version of the problem...

1. One ad per query
2. Every advertiser has the same budget
3. Every ad has the same click through rate
4. All bids are either 0 or 1

(either the advertiser wants the query, or they don't)

Greedy solution

Then the greedy solution is...

- Every time a new query comes in, select any advertiser who has bid on that query (who has budget remaining)
 - What is the competitive ratio of this algorithm?

Greedy solution

The balance algorithm

A better algorithm...

- Every time a new query comes in, amongst advertisers who have bid on this query, **select the one with the largest remaining budget**
 - How would this do on the same sequence?

The balance algorithm

A better algorithm...

- Every time a new query comes in, amongst advertisers who have bid on this query, **select the one with the largest remaining budget**
- In fact, the competitive ratio of this algorithm (still with equal budgets and fixed bids) is $(1 - 1/e) \sim 0.63$

The balance algorithm

What if bids aren't equal?

Bidder	Bid (on q)	Budget
A	1	110
B	10	100

The balance algorithm

What if bids aren't equal?

Bidder	Bid (on q)	Budget
A		
B		

The balance algorithm v2

We need to make two modifications

- We need to consider the bid amount when selecting the advertiser, and bias our selection toward higher bids
 - We also want to use some of each advertiser's budget (so that we don't just ignore advertisers whose budget is small)

The balance algorithm v2

Advertiser: A_i

fraction of budget remaining: f_i

bid on query q : $x_i(q)$

Assign queries to whichever advertiser maximizes:

$$\Psi_i(q) = x_i(q) \cdot (1 - e^{-f_i})$$

(could multiply by click-through rate if click-through rates are not equal)

The balance algorithm v2

Properties

- This algorithm has a competitive ratio of $(1 - \frac{1}{e})$.
- In fact, there **is no online algorithm** for the adwords problem with a competitive ratio **better than** $(1 - \frac{1}{e})$.

(proof is too deep for me...)

So far we have seen...

- An **online** algorithm to match advertisers to users (really to queries) that handles both **bids** and **budgets**
 - We wanted our **online** algorithm to be as good as the **offline** algorithm would be – we measured this using the **competitive ratio**
- Using a specific scheme that favored high bids while trying to balance the budgets of all advertisers, we achieved a ratio of $(1 - \frac{1}{e})$.
 - And no better online algorithm exists!

We **haven't** seen...

- AdWords actually uses a **second-price** auction (the winning advertiser pays the amount that the **second** highest bidder bid)
- Advertisers don't bid on specific queries, but inexact matches ('broad matching') – i.e., queries that include subsets, supersets, or synonyms of the keywords being bid on

Questions?

Further reading:

- Mining of Massive Datasets – “The Adwords Problem”
<http://infolab.stanford.edu/~ullman/mmds/book.pdf>
- AdWords and Generalized On-line Matching (A. Mehta)
<http://web.stanford.edu/~saber/adwords.pdf>