

Web Mining and Recommender Systems

Recommender Systems: Introduction

Why recommendation?

The goal of recommender systems is...


- To help people discover new content

Recommendations for You in Amazon Instant Video [See more](#)




Why recommendation?

- The goal of recommender systems is...
- To help us find the content we were already looking for



Why recommendation?

- The goal of recommender systems is...
- To discover which things go together




Customers Who Bought This Item Also Bought



Why recommendation?

The goal of recommender systems is...


- To personalize user experiences in response to user feedback



Why recommendation?

The goal of recommender systems is...

- To recommend incredible products that are relevant to our interests




Why recommendation?

The goal of recommender systems is...

- To identify things that we **like**

Results for 'mad max'



Add

★★★★☆

Mad Max
1979 **R** 93 minutes

In a postapocalyptic future, jaded motorcycle cop Max Rockatansky is ready to retire. But his world is shattered when a malicious gang murders his family as an act of retaliation, forcing a devastated Max to hit the open road seeking vengeance.

Starring: Mel Gibson, Hugh Keays-Byrne
Director: George Miller
Genre: Sci-Fi & Fantasy
Format: DVD

★★★★☆ **3.6** Our best guess for Jeremy

Why recommendation?

The goal of recommender systems is...

- To help people discover new content
- To help us find the content we were already looking for
- To discover which things go together
- To personalize user experiences in response to user feedback
 - To identify things that we **like**

Why recommendation?

The goal of recommender systems is...




- To help people discover new content
- To help us find the content we were
- To discover preferences, opinions, together
- To provide recommendations in response to user feedback
- To identify things that we **like**

To **model** people's preferences, opinions, and behavior

Recommending things to people

Suppose we want to build a movie recommender

e.g. which of these films will I rate highest?



Recommending things to people

We already have
a few tools in our
"supervised
learning" toolbox
that may help us

Pitch Black - Unrated Director's Cut [R] [CC]

★★★★★ 777 **IMDb** 7.1/10

Watch Trailer

When their ship crash-lands on a remote planet, the marooned passengers soon learn that escaped convict Riddick (Vin Diesel) isn't the only thing they have to fear. Deadly creatures lurk in the shadows, waiting to attack in the dark, and the planet is rapidly plunging into the

See More

Starring: Vin Diesel, Radha Mitchell
Runtime: 1 hour, 53 minutes
Available to watch on supported devices.

Product Details

Genres	Science Fiction, Action, Horror
Director	David Twohy
Starring	Vin Diesel, Radha Mitchell
Supporting actors	Cole Hauser, Keith David, Lewis Fitz-Gerald, Claudia Black, Rhiana Granger, Angela Moore, Peter Chiang, Ken Twohy
Studio	NBC Universal
MPAA Rating	R (Restricted)
Captions and subtitles	English Details
Rental rights	24 hour viewing period. Details
Purchase rights	Stream instantly and download to 2 locations Details
Format	Amazon Instant Video (streaming online video and digital download)

A. Phillips
Reviewer ranking: #17,230,554
90% helpful
votes received on reviews (151 of 167)

ABOUT ME
Enjoy the reviews...

ACTIVITIES
[Reviews \(16\)](#)
[Public Wish List \(2\)](#)
[Listmania Lists \(2\)](#)
[Tagged Items \(1\)](#)

$f(\text{user features, movie features}) \xrightarrow{?} \text{star rating}$

Recommending things to people

$f(\text{user features, movie features}) \xrightarrow{?} \text{star rating}$

User features: age, gender,
location, etc.

A. Phillips

Reviewer ranking: #17,230,554

90% helpful

votes received on reviews
(151 of 167)

ABOUT ME

Enjoy the reviews...

ACTIVITIES

[Reviews \(16\)](#)

[Public Wish List \(2\)](#)

[Listmania Lists \(2\)](#)

[Tagged Items \(1\)](#)

Movie features: genre,
actors, rating, length, etc.

Product Details

Genres	Science Fiction , Action , Horror
Director	David Twohy
Starring	Vin Diesel , Radha Mitchell
Supporting actors	Cole Hauser , Keith David , Lewis Fitz-Gerald , Claudia Black , Rhiana Gr Angela Moore , Peter Chiang , Ken Twohy
Studio	NBC Universal
MPAA rating	R (Restricted)
Captions and subtitles	English Details ▾
Rental rights	24 hour viewing period. Details ▾
Purchase rights	Stream instantly and download to 2 locations Details ▾
Format	Amazon Instant Video (streaming online video and digital download)

Recommending things to people

$f(\text{user features, movie features}) \xrightarrow{?} \text{star rating}$

With the models we've seen so far, we can build predictors that account for...

- Do women give higher ratings than men?
- Do Americans give higher ratings than Australians?
- Do people give higher ratings to action movies?
- Are ratings higher in the summer or winter?
- Do people give high ratings to movies with Vin Diesel?

So what **can't** we do yet?

Recommending things to people

$f(\text{user features}, \text{movie features}) \xrightarrow{?} \text{star rating}$

Consider the following linear predictor
(e.g. from week 1):

$$f(\text{user features}, \text{movie features}) = \langle \phi(\text{user features}); \phi(\text{movie features}), \theta \rangle$$

feature vector

$$\Rightarrow \langle \phi_1(\text{user}), \theta^{\text{user}} \rangle + \langle \phi_2(\text{movie}), \theta^{\text{movie}} \rangle$$

Recommending things to people

$f(\text{user features}, \text{movie features}) \xrightarrow{?} \text{star rating}$

Consider the following linear predictor
(e.g. from week 1):

$$\begin{aligned} f(\text{user features}, \text{movie features}) &= \\ &\quad \langle \phi(\text{user features}); \phi(\text{movie features}), \theta \rangle \\ &= \langle \phi(\text{user features}), \theta_{\text{user}} \rangle + \langle \phi(\text{movie features}), \theta_{\text{movie}} \rangle \end{aligned}$$


Recommending things to people

But this is essentially just two separate predictors!

$$\begin{aligned} f(\text{user features}, \text{movie features}) &= \\ &= \underbrace{\langle \phi(\text{user features}), \theta_{\text{user}} \rangle}_{\text{user predictor}} + \underbrace{\langle \phi(\text{movie features}), \theta_{\text{movie}} \rangle}_{\text{movie predictor}} \end{aligned}$$


That is, we're treating user and movie features as though they're **independent**!

Recommending things to people

But these predictors should (obviously?)
not be independent

$$f(\text{user features, movie features}) = f(\text{user}) + f(\text{movie})$$

do I tend to give high ratings?




does the population tend to give high ratings to this genre of movie?



But what about a feature like “do I give
high ratings to **this genre** of movie”?

Recommending things to people

Recommender Systems go beyond the methods we've seen so far by trying to model the **relationships** between people and the items they're evaluating



This section

Recommender Systems

1. (next) Collaborative filtering

(performs recommendation in terms of user/user and item/item similarity)

2. (later) Latent-factor models

(performs recommendation by projecting users and items into some low-dimensional space)

3. The Netflix Prize

4. Recommender Systems Evaluation

Later

Recommender Systems – more advanced topics

- Incorporating complex *side-information* into recommender systems
- Recommendation in other contexts, e.g. social networks, online dating, etc.
 - Online advertising
- (even later) temporal factors, ethics, text, etc.

Web Mining and Recommender Systems

Similarity-based Recommender Systems

Defining similarity between users & items

Q: How can we measure the **similarity** between two **users**?

A: In terms of the **items** they purchased!

Q: How can we measure the similarity between two **items**?

A: In terms of the users who purchased them!

Defining similarity between users & items

e.g.:
Amazon



Calvin Klein Men's Relaxed Straight Leg Jean In Cove

★★★★★ 20 customer reviews

Price: \$48.16 - \$69.99 & FREE Returns. Details

Size:

Select

Sizing info

Fit: As expected (55%)

Color: Cove

- 98% Cotton/2% Elastane
- Imported
- Button closure
- Machine Wash
- Relaxed straight-leg jean in light-tone denim featuring whiskering and five-pocket styling
- Zip fly with button
- 10.25-inch front rise, 19-inch knee, 17.5-inch leg opening

Frequently Bought Together



Calvin Klein Jeans
\$57.94 - \$69.50



Calvin Klein Jeans
\$49.92



Calvin Klein Jeans
\$50.67 - \$69.99



Levi's
\$23.99 - \$68.00

Customers Who Viewed This Item Also Viewed



Customers Who Bought This Item Also Bought



Page 2

Definitions

I_u = set of items purchased by user u

U_i = set of users who purchased item i


Definitions

Or equivalently...

$$R = \begin{pmatrix} 1 & 0 & \dots & 1 \\ 0 & 0 & & 1 \\ \vdots & & \ddots & \vdots \\ 1 & 0 & \dots & 1 \end{pmatrix}$$

items

users



R_u = binary representation of items purchased by u

$R_{\cdot, i}$ = binary representation of users who purchased i


$$I_u = \{i \mid R_{ui} = 1\} \quad U_i = \{u \mid R_{ui} = 1\}$$

0. Euclidean distance

Euclidean distance:

e.g. between two items i, j (similarly defined between two users)

$$|U_i \setminus U_j| + |U_j \setminus U_i| = \|R_i - R_j\|$$



0. Euclidean distance

Euclidean distance:

$$\begin{aligned}\text{e.g.: } U_1 &= \{1, 4, 8, 9, 11, 23, 25, 34\} \\ U_2 &= \{1, 4, 6, 8, 9, 11, 23, 25, 34, 35, 38\} \\ U_3 &= \{4\} \\ U_4 &= \{5\}\end{aligned}$$

$$|U_1 \setminus U_2| + |U_2 \setminus U_1| = 3$$

$$|U_3 \setminus U_4| + |U_4 \setminus U_3| = 2$$

Problem: favors small sets, even if they have few elements in common


1. Jaccard similarity

$$\text{Jaccard}(A, B) =$$

$$|A \cap B| / |A \cup B|$$

$$\text{Jaccard}(U_i, U_j) =$$


$$|u_i \cap u_j| / |u_i \cup u_j|$$



→ Maximum of 1 if the two users purchased **exactly the same** set of items
(or if two items were purchased by the same set of users)

→ Minimum of 0 if the two users purchased **completely disjoint** sets of items
(or if the two items were purchased by completely disjoint sets of users)

2. Cosine similarity



$$\cos(\theta) = 1$$

(theta = 0) $\rightarrow$ A and B point in exactly the same direction

$$\cos(\theta) = -1$$

(theta = 180) $\rightarrow$ A and B point in opposite directions (won't actually happen for 0/1 vectors)

$$\cos(\theta) = 0$$

(theta = 90) $\rightarrow$ A and B are orthogonal

$$\theta = \cos^{-1}\left(\frac{a \cdot b}{\|a\| \|b\|}\right) \quad \cos \theta = \frac{u_i \cdot u_j}{\|u_i\| \|u_j\|}$$

$$\text{for binary } = \frac{|u_i \cap u_j|}{\sqrt{|u_i|} \sqrt{|u_j|}}$$

2. Cosine similarity

Why cosine?

- Unlike Jaccard, works for arbitrary vectors
- E.g. what if we have **opinions** in addition to purchases?

$$R = \begin{pmatrix} 1 & 0 & \dots & 1 \\ 0 & 0 & & 1 \\ \vdots & & \ddots & \vdots \\ 1 & 0 & \dots & 1 \end{pmatrix} \longrightarrow \begin{pmatrix} -1 & 0 & \dots & 1 \\ 0 & 0 & & -1 \\ \vdots & & \ddots & \vdots \\ 1 & 0 & \dots & -1 \end{pmatrix}$$


bought and **liked**

didn't buy

bought and **hated**

2. Cosine similarity

E.g. our previous example, now with
"thumbs-up/thumbs-down" ratings



$$\cos(\theta) = 1$$

(theta = 0) -> Rated by the
same users, and they all agree

$$\cos(\theta) = -1$$

(theta = 180) -> Rated by the
same users, but they
completely disagree about it

$$\cos(\theta) = 0$$

(theta = 90) -> Rated by
different sets of users

3. Pearson correlation

What if we have numerical ratings
(rather than just thumbs-up/down)?

$$R = \begin{pmatrix} -1 & 0 & \dots & 1 \\ 0 & 0 & & -1 \\ \vdots & & \ddots & \vdots \\ 1 & 0 & \dots & -1 \end{pmatrix} \longrightarrow \begin{pmatrix} 4 & 0 & \dots & 2 \\ 0 & 0 & & 3 \\ \vdots & & \ddots & \vdots \\ 5 & 0 & \dots & 1 \end{pmatrix}$$



bought and **liked**

didn't buy

bought and **hated**

3. Pearson correlation

What if we have numerical ratings
(rather than just thumbs-up/down)?



3. Pearson correlation

What if we have numerical ratings
(rather than just thumbs-up/down)?

- We wouldn't want 1-star ratings to be parallel to 5-star ratings
- So we can subtract the average – values are then **negative** for below-average ratings and **positive** for above-average ratings

$$\text{Sim}(u, v) = \frac{\sum_{i \in I_u \cap I_v} (R_{u,i} - \bar{R}_u)(R_{v,i} - \bar{R}_v)}{\sqrt{\sum_{i \in I_u \cap I_v} (R_{u,i} - \bar{R}_u)^2 \sum_{i \in I_u \cap I_v} (R_{v,i} - \bar{R}_v)^2}}$$

items rated by both users average rating by user v

3. Pearson correlation

Compare to the cosine similarity:

Pearson similarity (between users):

$$\text{Sim}(u, v) = \frac{\sum_{i \in I_u \cap I_v} (R_{u,i} - \bar{R}_u)(R_{v,i} - \bar{R}_v)}{\sqrt{\sum_{i \in I_u \cap I_v} (R_{u,i} - \bar{R}_u)^2 \sum_{i \in I_u \cap I_v} (R_{v,i} - \bar{R}_v)^2}}$$

items rated by both users average rating by user v

Cosine similarity (between users):

$$\text{Sim}(u, v) = \frac{\sum_{i \in I_u \cap I_v} R_{u,i} R_{v,i}}{\sqrt{\sum_{i \in I_u \cap I_v} R_{u,i}^2 \sum_{i \in I_u \cap I_v} R_{v,i}^2}}$$

Note: slightly different from previous definition. Here similarity is determined only based on items *both* users have consumed

3. Pearson correlation

$$\text{Sim}(u, v) = \frac{\sum_{i \in I_u \cap I_v} R_{u,i} R_{v,i}}{\sqrt{\sum_{i \in I_u \cap I_v} R_{u,i}^2 \sum_{i \in I_u \cap I_v} R_{v,i}^2}}$$

$$\text{Cosine}(A, B) = \frac{A \cdot B}{\|A\| \|B\|}$$


(-1, 2, -1, 4)
(5, 2, 1, 1)

Consider **all items** in the denominator, or just shared items?

Just shared: two users should be considered maximally similar if they've rated shared items the same way. If only one user has rated an item, we have no evidence that the other user is different.

All: Two users who've rated items the same way *and only rated the same items* should be more similar than two users who've rated some different items.

Ultimately, these are *heuristics*, and either definition could be used depending on the situation

Collaborative filtering in practice

How does amazon generate their recommendations?

Given a product:



Let U_i be the set of users who viewed it

Rank products according to: $\frac{|U_i \cap U_j|}{|U_i \cup U_j|}$ (or cosine/pearson)



.86



.84



.82



.79



...



Collaborative filtering in practice

- Amazon uses the cosine similarity
- Similarity is defined between users: the goal is to recommend items that have previously been purchased by similar customers (e.g. "customers who bought items in your shopping cart also bought")
- Main challenges involve scalability: how to cluster users so that we can quickly identify similar users

Collaborative filtering in practice

Note: (surprisingly) that we built something pretty useful out of **nothing but interaction data** – we didn't look at any features of the products (or users!) whatsoever

Collaborative filtering in practice

But: we still have
a few problems left to address...

1. This is actually kind of slow given a huge enough dataset – if one user purchases one item, this will change the rankings of **every other item that was purchased by at least one user in common**
2. Of no use for **new users** and **new items** (“cold-start” problems)
3. Won’t necessarily encourage diverse results

Web Mining and Recommender Systems

Similarity based recommender – implementation

Code on course webpage

Uses Amazon "Musical Instrument" data from
<https://s3.amazonaws.com/amazon-reviews-pds/tsv/index.txt>

Code: Reading the data

Read the data:

```
In [1]: import gzip
        from collections import defaultdict
        import random
        import numpy
        import scipy.optimize
```

```
In [2]: path = "/home/jmcauley/datasets/mooc/amazon/amazon_reviews_us_Musical_Instruments/v1_00.tsv.gz"
```

```
In [3]: f = gzip.open(path, 'rt', encoding="utf8")
```

```
In [4]: header = f.readline()
        header = header.strip().split('\t')
```

Code: Reading the data

Our goal is to make recommendations of products based on users' purchase histories. The only information needed to do so is **user and item IDs**

```
In [5]: dataset = []
```

```
In [6]: for line in f:
        fields = line.strip().split('\t')
        d = dict(zip(header, fields))
        d['star_rating'] = int(d['star_rating'])
        d['helpful_votes'] = int(d['helpful_votes'])
        d['total_votes'] = int(d['total_votes'])
        dataset.append(d)
```

```
In [7]: dataset[0]
```


```
Out[7]: {'marketplace': 'US',
        'customer_id': '45610553',
        'review_id': 'RMDCHHHY5QZ9',
        'product_id': 'B00HH62VB6',
        'product_parent': '618218723',
        'product_title': 'AGPtek® 10 Isolated Output 9V 12V 18V Guitar Pedal Board Power Supply Effect Pedals with Isolated Short Cricuit / Overcurrent Protection',
```

Code: Useful data structures

Build data structures representing the set of items for each user and users for each item:

```
In [8]: # Useful data structures
```

```
In [9]: usersPerItem = defaultdict(set)  
        itemsPerUser = defaultdict(set)
```



```
In [10]: itemNames = {}
```

```
In [11]: for d in dataset:  
        user,item = d['customer_id'], d['product_id']  
        usersPerItem[item].add(user)  
        itemsPerUser[user].add(item)  
        itemNames[item] = d['product_title']
```

Code: Jaccard similarity

The Jaccard similarity implementation follows the definition directly:

$$\text{Jaccard}(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

```
In [12]: def Jaccard(s1, s2):  
         numer = len(s1.intersection(s2))  
         denom = len(s1.union(s2))  
         return numer / denom
```

Recommendation

We want a recommendation function that return **items similar to a candidate item i** . Our strategy will be as follows:

- Find the set of users who purchased i
- Iterate over all other items other than i
- For all other items, compute their similarity with i
(*and store it*)
- Sort all other items by (Jaccard) similarity
 - Return the most similar

Code: Recommendation

Now we can implement the recommendation function itself:

```
In [13]: def mostSimilar(i):  
    similarities = []  
    users = usersPerItem[i]  
    for i2 in usersPerItem:  
        if i2 == i: continue  
        sim = Jaccard(users, usersPerItem[i2])  
        similarities.append((sim,i2))  
    similarities.sort(reverse=True)  
    return similarities[:10]
```

$$\text{Jaccard}(U_i, U_j) = \frac{|U_i \cap U_j|}{|U_i \cup U_j|}$$



Code: Recommendation

Next, let's use the code to make a recommendation.
The query is just a product ID:

```
In [14]: dataset[2]
```

```
Out[14]: {'marketplace': 'US',  
          'customer_id': '6111003',  
          'review_id': 'RIZR67JKUDBI0',  
          'product_id': 'B0006VMBHI',  
          'product_parent': '603261968',  
          'product_title': 'AudioQuest LP record clean brush',  
          'product_category': 'Musical Instruments',  
          'star_rating': 3,  
          'helpful_votes': 0,  
          'total_votes': 1,  
          'vine': 'N',  
          'verified_purchase': 'Y',  
          'review_headline': 'Three Stars',  
          'review_body': 'removes dust. does not clean',  
          'review_date': '2015-08-31'}
```

```
In [15]: query = dataset[2]['product_id']
```

Code: Recommendation

Next, let's use the code to make a recommendation.
The query is just a product ID:

```
In [16]: mostSimilar(query)
```

```
Out[16]: [(0.028446389496717725, 'B00006I5SD'),  
(0.01694915254237288, 'B00006I5SB'),  
(0.015065913370998116, 'B000AJR482'),  
(0.014204545454545454, 'B00E7MVP3S'),  
(0.008955223880597015, 'B001255YL2'),  
(0.008849557522123894, 'B003EIRV08'),  
(0.008333333333333333, 'B0015VEZ22'),  
(0.00821917808219178, 'B00006I5UH'),  
(0.008021390374331552, 'B00008BWM7'),  
(0.007656967840735069, 'B000H2BC4E')]
```

Code: Recommendation

Items that were recommended:

```
In [17]: itemNames[query]
```

```
Out[17]: 'AudioQuest LP record clean brush'
```

```
In [18]: [itemNames[x[1]] for x in mostSimilar(query)]
```

```
Out[18]: ['Shure SFG-2 Stylus Tracking Force Gauge',  
          'Shure M97xE High-Performance Magnetic Phono Cartridge',  
          'ART Pro Audio DJPRE II Phono Turntable Preamplifier',  
          'Signstek Blue LCD Backlight Digital Long-Playing LP Turntable Stylus Force Scale Gauge Tester',  
          'Audio Technica AT120E/T Standard Mount Phono Cartridge',  
          'Technics: 45 Adaptor for Technics 1200 (SFWE010)',  
          'GruvGlide GRUVGLIDE DJ Package',  
          'STANTON MAGNETICS Record Cleaner Kit',  
          'Shure M97xE High-Performance Magnetic Phono Cartridge',  
          'Behringer PP400 Ultra Compact Phono Preamplifier']
```

Recommending more efficiently

Our implementation was not very efficient. The slowest component is the iteration over all other items:

- Find the set of users who purchased i
- **Iterate over all other items other than i**
- For all other items, compute their similarity with i
(*and store it*)
- Sort all other items by (Jaccard) similarity
 - Return the most similar

This can be done more efficiently as most items will have no overlap

Recommending more efficiently

In fact it is sufficient to iterate over **those items purchased by one of the users who purchased i**

- Find the set of users who purchased i
- **Iterate over all users who purchased i**
- Build a candidate set from all items those users consumed
- For items in this set, compute their similarity with i (*and store it*)
 - Sort all other items by (Jaccard) similarity
 - Return the most similar

Code: Faster implementation

Our more efficient implementation works as follows:

```
In [19]: def mostSimilarFast(i):  
    similarities = []  
    users = usersPerItem[i]  
    candidateItems = set()  
    for u in users:  
        candidateItems = candidateItems.union(itemsPerUser[u])  
    for i2 in candidateItems:  
        if i2 == i: continue  
        sim = Jaccard(users, usersPerItem[i2])  
        similarities.append((sim,i2))  
    similarities.sort(reverse=True)  
    return similarities[:10]
```

Code: Faster recommendation

Which ought to recommend the same set of items, but
much more quickly:

```
In [20]: mostSimilarFast(query)
```

```
Out[20]: [(0.028446389496717725, 'B00006I5SD'),  
          (0.01694915254237288, 'B00006I5SB'),  
          (0.015065913370998116, 'B000AJR482'),  
          (0.014204545454545454, 'B00E7MVP3S'),  
          (0.008955223880597015, 'B001255YL2'),  
          (0.008849557522123894, 'B003EIRV08'),  
          (0.008333333333333333, 'B0015VEZ22'),  
          (0.00821917808219178, 'B00006I5UH'),  
          (0.008021390374331552, 'B00008BWM7'),  
          (0.007656967840735069, 'B000H2BC4E')]
```

Web Mining and Recommender Systems

Similarity-based rating prediction

Collaborative filtering for rating prediction

In the previous section we provided code to make recommendations based on the **Jaccard similarity**

How can the same ideas be used for rating prediction?

Collaborative filtering for rating prediction

A simple heuristic for rating prediction works as follows:

- The user (u)'s rating for an item i is a weighted combination of all of their previous ratings for items j
- The weight for each rating is given by the Jaccard similarity between i and j

Collaborative filtering for rating prediction

This can be written as:

$$r(u, i) = \frac{1}{Z} \sum_{j \in I_u \setminus \{i\}} r_{uj} \times \underbrace{\text{Sim}(i, j)}_{\substack{\text{similarity} \\ \text{(Jaccard etc.)}}}$$

$Z = \sum_{j \in I_u \setminus \{i\}} \text{Sim}(i, j)$

$j \in I_u \setminus \{i\}$
all other items rated by u

Collaborative filtering for rating prediction

This can be written as:

$$r(u, i) = \frac{1}{Z} \sum_{j \in I_u \setminus \{i\}} r_{u,j} \cdot \text{sim}(i, j)$$

Normalization
constant

All items the user has
rated other than i

$$Z = \sum_{j \in I_u \setminus \{i\}} \text{sim}(i, j)$$

Collaborative filtering for rating prediction

Other rating prediction functions...

$$r(u, i) = \frac{1}{Z} \sum_{v \in U_i \setminus \{u\}} r_{v,i} \text{Sim}(u, v)$$

$$r(u, i) = \underbrace{\bar{R}_i}_{\text{average rating for item } i} + \frac{1}{Z} \sum_{j \in I_u \setminus \{i\}} (r_{u,j} - \bar{R}_j) \text{Sim}(i, j)$$

Code: CF for rating prediction

Now we can adapt our previous recommendation code to predict ratings

```
In [22]: # More utility data structures
```

```
In [23]: reviewsPerUser = defaultdict(list)
reviewsPerItem = defaultdict(list)
```

List of reviews per user and per item

```
In [24]: for d in dataset:
        user,item = d['customer_id'], d['product_id']
        reviewsPerUser[user].append(d)
        reviewsPerItem[item].append(d)
```

```
In [25]: ratingMean = sum([d['star_rating'] for d in dataset]) / len(dataset)
```

```
In [26]: ratingMean
```

```
Out[26]: 4.251102772543146
```


We'll use the mean rating as a baseline for comparison

Code: CF for rating prediction

Our rating prediction code works as follows:

In [27]: `def predictRating(user,item):`

```
    ratings = []
    similarities = []
    for d in reviewsPerUser[user]:
        i2 = d['product_id']
        if i2 == item: continue
        ratings.append(d['star_rating'])
        similarities.append(Jaccard(usersPerItem[item],usersPerItem[i2]))
    if (sum(similarities) > 0):
        weightedRatings = [(x*y) for x,y in zip(ratings,similarities)]
        return sum(weightedRatings) / sum(similarities)
    else:
        # User hasn't rated any similar items
        return ratingMean
```

$$r(u, i) = \frac{1}{Z} \sum_{j \in I_u \setminus \{i\}} r_{u,j} \cdot \text{sim}(i, j)$$


Code: CF for rating prediction

As an example, select a rating for prediction:

```
In [28]: dataset[1]
```

```
Out[28]: {'marketplace': 'US',  
         'customer_id': '14640079',  
         'review_id': 'RZSL0BALIYUNU',  
         'product_id': 'B003LRN53I',  
         'product_parent': '986692292',  
         'product_title': 'Sennheiser HD203 Closed-Back DJ Headphones',  
         'product_category': 'Musical Instruments',  
         'star_rating': 5,  
         'helpful_votes': 0,  
         'total_votes': 0,  
         'vine': 'N',  
         'verified_purchase': 'Y',  
         'review_headline': 'Five Stars',  
         'review_body': 'Nice headphones at a reasonable price.',  
         'review_date': '2015-08-31'}
```

```
In [29]: u,i = dataset[1]['customer_id'], dataset[1]['product_id']
```

```
In [30]: predictRating(u, i)
```

```
Out[30]: 5.0
```

Code: CF for rating prediction

Similarly, we can evaluate accuracy across the entire corpus:

```
In [31]: def MSE(predictions, labels):  
         differences = [(x-y)**2 for x,y in zip(predictions,labels)]  
         return sum(differences) / len(differences)
```

```
In [32]: alwaysPredictMean = [ratingMean for d in dataset]
```

```
In [33]: cfPredictions = [predictRating(d['customer_id'], d['product_id']) for d in dataset]
```

```
In [34]: labels = [d['star_rating'] for d in dataset]
```

```
In [35]: MSE(alwaysPredictMean, labels)
```

```
Out[35]: 1.4796142779564334
```

```
In [36]: MSE(cfPredictions, labels)
```

```
Out[36]: 1.6146130004291603
```

Collaborative filtering for rating prediction

Note that this is just a **heuristic** for rating prediction

- In fact in this case it did *worse* (in terms of the MSE) than always predicting the mean
 - We could adapt this to use:
 1. A different similarity function (e.g. cosine)
 2. Similarity based on users rather than items
 3. A different weighting scheme

Better heuristics?

- interchange users / items
- normalize differently
- different $\text{Sim}()$ (Jacc./Pearson/Cosine etc)
- data other than interactions (eg features/time)

Web Mining and Recommender Systems

Latent-factor models

Summary so far

Recap

1. Measuring similarity between users/items for **binary** prediction
Jaccard similarity
2. Measuring similarity between users/items for **real-valued** prediction
cosine/Pearson similarity

Now: Machine learning-based models for **real-valued** prediction *latent-factor models*

Latent factor models

So far we've looked at approaches that try to define some definition of user/user and item/item **similarity**

Recommendation then consists of

- Finding an item i that a user likes (gives a high rating)
- Recommending items that are similar to it (i.e., items j with a similar rating profile to i)

Latent factor models

What we've seen so far are
unsupervised approaches and whether
the work depends highly on whether we
chose a "good" notion of similarity

So, can we perform recommendations
via **supervised** learning?

Latent factor models

e.g. if we can model

$f(\text{user features, movie features}) \rightarrow \text{star rating}$

Then recommendation
will consist of identifying

$$\textit{recommendation}(u) = \arg \max_{i \in \text{unseen items}} f(u, i)$$

The Netflix prize

In 2006, Netflix created a dataset of **100,000,000** movie ratings

Data looked like:

(userID, itemID, time, rating)

The goal was to reduce the (R)MSE at predicting ratings:

$$\text{RMSE}(f) = \sqrt{\frac{1}{N} \sum_{u,i,t \in \text{test set}} (f(u,i,t) - r_{u,i,t})^2}$$


model's prediction ground-truth

Whoever first manages to reduce the RMSE by **10%** versus
Netflix's solution wins **\$1,000,000**

The Netflix prize

This led to **a lot** of research on rating prediction by minimizing the Mean-Squared Error

NETFLIX

(it also led to a lawsuit against Netflix, once somebody managed to de-anonymize their data)

We'll look at a few of the main approaches

Rating prediction

Let's start with the simplest possible model:

$$f(u, i) = \alpha$$

↑ ↑
user item

$$\alpha = \bar{r}$$

θ_0

Rating prediction

What about the **2nd** simplest model?

$$f(u, i) = \alpha + \beta_u + \beta_i$$

user item

how much does
this user tend to
rate things above
the mean?

does this item tend
to receive higher
ratings than others

e.g.

$$\alpha = 4.2$$



$$\beta_{\text{pitch black}} = -0.1$$

$$\beta_{\text{julian}} = -0.2$$



Rating prediction

$$f(u, i) = \alpha + \beta_u + \beta_i$$

This is a linear model!

$$X(u, i) = [1, \underset{\substack{\uparrow \\ \text{user}}}{0 \ 0 \ 0 \ 0 \ 0}, \underset{\substack{\uparrow \\ \text{item}}}{0 \ 0 \ 0 \ 1 \ 0 \ 0 \ 0 \ 0}]$$
$$\Theta = [\alpha, \underset{\substack{\uparrow \\ \beta^{(\text{user})}}}{}, \underset{\substack{\uparrow \\ \beta^{(\text{item})}}}{}]$$


Rating prediction

The optimization problem becomes:

$$\arg \min_{\alpha, \beta} \underbrace{\sum_{u,i} (\alpha + \beta_u + \beta_i - R_{u,i})^2}_{\text{error}} + \lambda \underbrace{[\sum_u \beta_u^2 + \sum_i \beta_i^2]}_{\text{regularizer}}$$

Jointly convex in β_i, β_u . Can be solved by iteratively removing the mean and solving for beta

Jointly convex?



Rating prediction

Differentiate:

$$\arg \min_{\alpha, \beta} \sum_{u,i} (\alpha + \beta_u + \beta_i - R_{u,i})^2 + \lambda [\sum_u \beta_u^2 + \sum_i \beta_i^2]$$

$$\frac{\partial \text{obj}}{\partial \beta_u} = 2 \sum_{i \in I_u} (\alpha + \beta_u + \beta_i - R_{u,i}) + 2\lambda \beta_u$$

Rating prediction



Differentiate:

$$\frac{\partial \text{obj}}{\partial \beta_u} = \sum_{i \in I_u} 2(\alpha + \beta_u + \beta_i - R_{u,i}) + 2\lambda\beta_u$$

Two ways to solve:

1. "Regular" gradient descent
2. Solve $\frac{\partial \text{obj}}{\partial \beta_u} = 0$ (sim. for β_i , α)

Rating prediction

Differentiate:

$$\frac{\partial \text{obj}}{\partial \beta_u} = \sum_{i \in I_u} 2(\alpha + \beta_u + \beta_i - R_{u,i}) + 2\lambda\beta_u$$

Solve $\frac{\partial \text{obj}}{\partial \beta_u} = 0$:

$$-\lambda\beta_u = \sum_{i \in I_u} (\alpha + \beta_u + \beta_i - R_{u,i})$$
$$\beta_u = \frac{\sum_{i \in I_u} (\alpha + \beta_i - R_{u,i})}{(-\lambda - |I_u|)}$$

Rating prediction

Iterative procedure – repeat the following updates until convergence:

$$\alpha = \frac{\sum_{u,i \in \text{train}} (R_{u,i} - (\beta_u + \beta_i))}{N_{\text{train}}}$$

$$\beta_u = \frac{\sum_{i \in I_u} R_{u,i} - (\alpha + \beta_i)}{\lambda + |I_u|}$$

$$\beta_i = \frac{\sum_{u \in U_i} R_{u,i} - (\alpha + \beta_u)}{\lambda + |U_i|}$$

(exercise: write down derivatives and convince yourself of these update equations!)

Rating prediction

Looks good (and actually works surprisingly well), but doesn't solve the basic issue that we started with

$$\begin{aligned} f(\text{user features}, \text{movie features}) &= \\ &= \underbrace{\langle \phi(\text{user features}), \theta_{\text{user}} \rangle}_{\text{user predictor}} + \underbrace{\langle \phi(\text{movie features}), \theta_{\text{movie}} \rangle}_{\text{movie predictor}} \end{aligned}$$


That is, we're **still** fitting a function that treats users and items independently

Web Mining and Recommender Systems

Latent-factor models (part 2)

Recommending things to people

How about an approach based on
dimensionality reduction?



i.e., let's come up with low-dimensional representations of the users and the items so as to best explain the data

Dimensionality reduction

~~We already have some tools that ought to help us, e.g. from dimensionality reduction:~~

if you've taken linear algebra!

$$R = \begin{pmatrix} 5 & 3 & \cdots & 1 \\ 4 & 2 & & 1 \\ 3 & 1 & & 3 \\ 2 & 2 & & 4 \\ 1 & 5 & & 2 \\ \vdots & & \ddots & \vdots \\ 1 & 2 & \cdots & 1 \end{pmatrix}$$

users

items

What is the best low-rank approximation of R in terms of the mean-squared error?

Dimensionality reduction

We can borrow some existing tools, e.g. the singular value decomposition, PCA (etc.):

$$R = \begin{pmatrix} 5 & 3 & \cdots & 1 \\ 4 & 2 & & 1 \\ 3 & 1 & & 3 \\ \vdots & & \ddots & \vdots \\ 1 & 2 & \cdots & 1 \end{pmatrix}$$

Singular Value Decomposition

(square roots of)
eigenvalues of RR^T

$$R = U \Sigma V^T$$

eigenvectors of RR^T

eigenvectors of $R^T R$

The “best” rank-K approximation (in terms of the MSE) consists of taking the eigenvectors with the highest eigenvalues

Dimensionality reduction

But! Our matrix of ratings is only partially observed; and it's **really big!**

$$R = \begin{pmatrix} 5 & 3 & \dots & \cdot \\ 4 & 2 & & 1 \\ 3 & \cdot & & 3 \\ \cdot & 2 & & 4 \\ 1 & 5 & & \cdot \\ \vdots & & \ddots & \vdots \\ 1 & 2 & \dots & \cdot \end{pmatrix}$$

Missing ratings

SVD is **not defined** for partially observed matrices, and it is **not practical** for matrices with 1Mx1M+ dimensions

Latent-factor models

Instead, let's solve approximately using
gradient descent

$$R = \underbrace{\begin{pmatrix} 5 & 3 & \dots & \cdot \\ 4 & 2 & & 1 \\ 3 & \cdot & & 3 \\ \cdot & 2 & & 4 \\ 1 & 5 & & \cdot \\ \vdots & & \ddots & \vdots \\ 1 & 2 & \dots & \cdot \end{pmatrix}}_{\text{items}} \left. \vphantom{\begin{pmatrix} 5 \\ 4 \\ 3 \\ \cdot \\ 1 \\ \vdots \\ 1 \end{pmatrix}} \right\} \text{users}$$

$R \simeq UV^T$

K-dimensional representation of each item

K-dimensional representation of each user

Latent-factor models

Instead, let's solve approximately using
gradient descent

$$R = \begin{pmatrix} 5 & 3 & \dots & \cdot \\ 4 & 2 & & 1 \\ 3 & \cdot & & 3 \\ \cdot & 2 & & 4 \\ 1 & 5 & & \cdot \\ \vdots & & \ddots & \vdots \\ 1 & 2 & \dots & \cdot \end{pmatrix} \approx \begin{bmatrix} \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \end{bmatrix} \begin{bmatrix} \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \\ \vdots \end{bmatrix}$$


Handwritten notes in red:

- γ_u (next to the first column vector)
- γ_i (next to the second column vector)
- $R_{ui} \approx \gamma_u \cdot \gamma_i$ (at the bottom right)
- γ_I (circled, next to the second column vector)

Latent-factor models

Let's write this as:

$$f(u, i) = \alpha + \beta_u + \beta_i + \gamma_u \cdot \gamma_i$$



Latent-factor models

Let's write this as:

$$f(u, i) = \alpha + \beta_u + \beta_i + \gamma_u \cdot \gamma_i$$

Our optimization problem is then

$$\arg \min_{\alpha, \beta, \gamma} \underbrace{\sum_{u,i} (\alpha + \beta_u + \beta_i + \gamma_u \cdot \gamma_i - R_{u,i})^2}_{\text{error}} + \lambda \underbrace{[\sum_u \beta_u^2 + \sum_i \beta_i^2 + \sum_i \|\gamma_i\|_2^2 + \sum_u \|\gamma_u\|_2^2]}_{\text{regularizer}}$$

Latent-factor models

Problem: this is certainly not convex

- ① It's smooth
- ② It has many local minima



Latent-factor models

Oh well. We'll just solve it approximately
Again, two ways to solve:

1. "Regular" gradient descent
2. Solve $\frac{\partial \text{obj}}{\partial \gamma_u} = 0$ (sim. For β_i , α , etc.)

(**Solution 1** is much easier to implement,
though **Solution 2** might converge more
quickly/easily)

Latent-factor models (Solution 1)

$$\arg \min_{\alpha, \beta, \gamma} \sum_{u,i} (\alpha + \beta_u + \beta_i + \gamma_u \cdot \gamma_i - R_{u,i})^2 + \lambda [\sum_u \beta_u^2 + \sum_i \beta_i^2 + \sum_i \|\gamma_i\|_2^2 + \sum_u \|\gamma_u\|_2^2]$$

$$\gamma = \sum_k \gamma_{uk} \cdot \gamma_{ik}$$

$$\frac{\partial \text{obj}}{\partial \gamma_{uk}} =$$

$$2 \sum_{i \in I_u} \gamma_{ik} (\alpha + \beta_u + \beta_i + \gamma_u \cdot \gamma_i - R_{u,i}) + 2\lambda \gamma_{uk}$$

Latent-factor models (Solution 2)

Observation: if we know either the user or the item parameters, the problem becomes "easy"

$$f(u, i) = \alpha + \beta_u + \beta_i + \gamma_u \cdot \gamma_i$$

e.g. fix γ_i – pretend we're fitting parameters for features



Latent-factor models

(Harder solution): iteratively solve the following subproblems

objective:

$$\arg \min_{\alpha, \beta, \gamma} \sum_{u,i} (\alpha + \beta_u + \beta_i + \gamma_u \cdot \gamma_i - R_{u,i})^2 + \lambda [\sum_u \beta_u^2 + \sum_i \beta_i^2 + \sum_i \|\gamma_i\|_2^2 + \sum_u \|\gamma_u\|_2^2]$$

$$= \arg \min_{\alpha, \beta, \gamma} \text{objective}(\alpha, \beta, \gamma)$$

1) fix γ_i . Solve $\arg \min_{\alpha, \beta, \gamma_u} \text{objective}(\alpha, \beta, \gamma)$

2) fix γ_u . Solve $\arg \min_{\alpha, \beta, \gamma_i} \text{objective}(\alpha, \beta, \gamma)$

3,4,5...) repeat until convergence

Each of these subproblems is “easy” – just regularized least-squares, like we’ve been doing since we studied regression.

This procedure is called **alternating least squares**.

Latent-factor models

later we'll see how to do this using:

- High-level recommender systems libraries
 - Tensorflow (maybe? I always run out of time so probably not)

Latent-factor models

Observation: we went from a method which uses **only** features:

$f(\text{user features, movie features}) \rightarrow \text{star rating}$

Diagram illustrating the features used in the initial method:

- User features:** age, gender, location, etc. (Screenshot of a user profile for A. Phillips).
- Movie features:** genre, actors, rating, length, etc. (Screenshot of a movie product details page).

to one which **completely ignores** them:

$$\arg \min_{\alpha, \beta, \gamma} \sum_{u,i} (\alpha + \beta_u + \beta_i + \gamma_u \cdot \gamma_i - R_{u,i})^2 + \lambda [\sum_u \beta_u^2 + \sum_i \beta_i^2 + \sum_i \|\gamma_i\|_2^2 + \sum_u \|\gamma_u\|_2^2]$$

Latent-factor models

Should we use features or not?

1) Argument **against** features:

In principle, the addition of features adds **no expressive power** to the model. We **could** have a feature like “is this an action movie?”, but if this feature were useful, the model would “discover” a latent dimension corresponding to action movies, and we wouldn’t need the feature anyway

In the limit, this argument is valid: as we add more ratings per user, and more ratings per item, the latent-factor model should automatically discover any useful dimensions of variation, so the influence of observed features will disappear

Latent-factor models

Should we use features or not?

2) Argument **for** features:

But! Sometimes we **don't** have many ratings per user/item

Latent-factor models are next-to-useless if **either** the user or the item was never observed before

$$f(u, i) = \alpha + \beta_u + \beta_i + \gamma_u \cdot \gamma_i$$

reverts to zero if we've never seen the user before
(because of the regularizer)

Latent-factor models

Should we use features or not?


2) Argument **for** features:

This is known as the **cold-start** problem in recommender systems. Features are not useful if we have many observations about users/items, but are useful for **new** users and items.

We also need some way to handle users who are **active**, but don't necessarily rate anything, e.g. through **implicit feedback**

Dimensionality reduction

Note that this is really a form of ***dimensionality reduction***



- What are the dimensions that *explain the most variance* in the data?
- For connections to other dimensionality reduction techniques (mostly SVD), see textbook

Overview & recap

Recently we've followed the programme below:

1. Measuring similarity between users/items for **binary** prediction (e.g. Jaccard similarity)
2. Measuring similarity between users/items for **real-valued** prediction (e.g. cosine/Pearson similarity)
3. Dimensionality reduction for **real-valued** prediction (latent-factor models)
4. **Finally** – dimensionality reduction for **binary** prediction

$\gamma_u \cdot \gamma_i$


Web Mining and Recommender Systems

Implicit feedback models

One-class recommendation

Suppose we have binary (0/1) observations
(e.g. purchases) or pos./neg. feedback
(thumbs-up/down)

$$R = \begin{pmatrix} 1 & 0 & \cdots & 1 \\ 0 & 0 & & 1 \\ \vdots & & \ddots & \vdots \\ 1 & 0 & \cdots & 1 \end{pmatrix} \text{ or } \begin{pmatrix} -1 & ? & \cdots & 1 \\ ? & ? & & -1 \\ \vdots & & \ddots & \vdots \\ 1 & ? & \cdots & -1 \end{pmatrix}$$



purchased didn't purchase liked didn't evaluate didn't like

One-class recommendation

How can we use **dimensionality reduction** (latent factors) to predict **binary** outcomes?

$x \cdot \theta$

$\sigma(x \cdot \theta)$

- Previously we saw **regression** and **logistic regression**.
These two approaches use the same type of linear function to predict real-valued and binary outputs
 - We can apply an analogous approach to binary recommendation tasks

This is referred to as “**one-class**” recommendation

Why can't we just apply logistic regression?

Why do we need a special approach? Compare to "traditional" approach of replacing "missing values" by 0:

	i_1	i_2	i_3	i_4	
u_1	?	+	+	?	user ↑ ↓
u_2	+	?	?	+	
u_3	+	+	?	?	
u_4	?	?	+	+	
u_5	?	?	+	?	
	← item →				



0	1	1	0
1	0	0	1
1	1	0	0
0	0	1	1
0	0	1	0

fit
 $\sigma(\gamma_u \cdot \gamma_i)$

Why can't we just apply logistic regression?

- 1) Data very imbalanced
- 2) (usually) can't deal with all zeros
- 3) Negatives aren't really negative

Why can't we just apply logistic regression?

Why do we need a special approach? Compare to “traditional” approach of replacing “missing values” by 0:

- At test time, the model should assign positive scores to items that the user consumed
 - But at training time, the model was penalized for *not* predicting zero!
- (Put differently, the “negative” items are exactly the ones we should be recommending!)

One-class recommendation

Two broad classes of strategy to dealing with one-class data:

1. Instance reweighting: try to figure out which negative (or positive) instances are "important"
2. Optimize *relative* scores rather than positive versus negative

Why can't we just apply logistic regression?

We need a special way to handle "negative" items (since they're not really "negative")

1. Try to figure out which negatives are "real" negatives, and weight instances differently (*instance reweighting*)
2. Try to use a ranking-based objective (*personalized ranking*)

Instance reweighting

1. Instance reweighting: try to figure out which negative (or positive) instances are "important"

Fit a function of the form:

$$\sum_{u_i} c_{u_i} (p_{u_i} - \underbrace{\sigma_u \cdot \sigma_i}_{\text{model } f(u,i)})^2$$

instance weight $label \in \{0, 1\}$

Instance reweighting

1. Instance reweighting: try to figure out which negative (or positive) instances are "important"

Fit a function of the form:

$$\arg \min_{\gamma} \sum_{(u,i) \in \mathcal{T}} c_{u,i} (p_{u,i} - \gamma_u \cdot \gamma_i)^2 + \lambda \Omega(\gamma)$$

could do: $c_{u,i} \sigma(p_{u,i} - \gamma_u \cdot \gamma_i)$

Instance reweighting

How to choose c (i.e., the importance of each sample)? A couple of heuristics:

1. (Hu et al. 2008): applied to positive instances

$$c_{u,i} = 1 + \alpha r_{u,i}$$

$$c_{u,i} = 1 + \alpha \log(1 + r_{u,i}/\epsilon)$$


$r_{u,i}$ # of times song was listened to

Instance reweighting

How to choose c (i.e., the importance of each sample)? A couple of heuristics:

2. (Pan et al. 2008): applied to negative instances

$$c_{u,i} = \alpha \times |I_u| \qquad c_{u,i} = \alpha(m - |U_i|)$$



(negative instances should be weighted higher if the user has interacted with many items, etc.)

Instance reweighting

- heuristics!
 - quick ways to incorporate side-information
- depend on semantics of a specific dataset

2. Bayesian Personalized Ranking

Idea: Rather than predicting that negative items are *disliked*, can we just predict that they're *less liked* than positive items?

$u:$ $++ ++ ? ? ? ? ? ? ? ?$

$\underbrace{\quad\quad\quad} \quad \underbrace{\quad\quad\quad} \quad \underbrace{\quad\quad\quad}$

$\sigma_u \cdot \sigma_i > \sigma_u \cdot \sigma_i > \sigma_u \cdot \sigma_i$

high ? low

Bayesian Personalized Ranking

Goal: Estimate a personalized **ranking function** for each user

- Compare pairs of items i and j together
- i is an item u consumed ("positive")
- j is an item u didn't consume
- Train such that i should have a higher score than j (for u)

Bayesian Personalized Ranking

Basic scheme:

- Our original dataset consists of *positives* (u,i) , e.g. purchased items for each user
- Augment this dataset by constructing many *triples* (u,i,j) where (u,i) is positive and (u,j) is negative
- The model now has to make binary predictions as to whether i or j is the positive item

Bayesian Personalized Ranking

Goal: Estimate a personalized **ranking function** for each user

$$x(u, i, j) > 0 \quad \text{if } i \text{ preferred}$$
$$x(u, i, j) < 0 \quad \text{otherwise}$$

Bayesian Personalized Ranking

What form should $x(u,i,j)$ take?


$$x(u,i,j) = f(u,i) - f(u,j)$$
$$= \sigma_u \cdot \sigma_i - \sigma_u \cdot \sigma_j$$

rec: $\arg\max_i f(u,i)$

Bayesian Personalized Ranking

Goal is to count how many times we identified i as being "more preferable" than j for a user u

$$\delta(\hat{x}_{uij} > 0)$$



Bayesian Personalized Ranking

We can think of this as maximizing the probability of correctly predicting pairwise preferences, i.e.,

$$p(i \text{ is preferred over } j) = \sigma(\gamma_u \cdot \gamma_i - \gamma_u \cdot \gamma_j)$$

- As with logistic regression, we can now maximize the likelihood associated with such a model by gradient ascent
 - In practice it isn't feasible to consider all pairs of positive/negative items, so we proceed by stochastic gradient ascent – i.e., randomly sample a (positive, negative) pair and update the model according to the gradient w.r.t. that pair

Bayesian Personalized Ranking

$$\max \ln \sigma(\gamma_u \cdot \gamma_i - \gamma_u \cdot \gamma_j)$$

$$\text{obj} = -\ln \left(1 - e^{\gamma_u \cdot \gamma_j - \gamma_u \cdot \gamma_i} \right)$$

$$\frac{\partial \text{obj}}{\partial \gamma_{uk}} = \sum_{i,j \in I_u} \frac{(\gamma_{ik} - \gamma_{jk}) e^{\gamma_u \cdot \gamma_j - \gamma_u \cdot \gamma_i}}{(1 + e^{\gamma_u \cdot \gamma_j - \gamma_u \cdot \gamma_i})}$$

Recap

1. Measuring similarity between users/items for **binary** prediction
Jaccard similarity
2. Measuring similarity between users/items for **real-valued** prediction
cosine/Pearson similarity
3. Dimensionality reduction for **real-valued** prediction
latent-factor models
4. Dimensionality reduction for **binary** prediction
one-class recommender systems

References

Further reading:

One-class recommendation:

<http://goo.gl/08Rh59>

Amazon's solution to collaborative filtering at scale:

<http://www.cs.umd.edu/~samir/498/Amazon-Recommendations.pdf>

An (expensive) textbook about recommender systems:

<http://www.springer.com/computer/ai/book/978-0-387-85819-7>

Cold-start recommendation (e.g.):

<http://wanlab.poly.edu/recsys12/recsys/p115.pdf>

Web Mining and Recommender Systems

Recommender Systems Libraries

Latent Factor Models (Surprise)

Fairly simple interface that implements the type of (rating prediction) model we've described

Reads data in .tsv format (various others are supported):


```
reader = Reader(line_format='user item rating', sep='\t')
data = Dataset.load_from_file(dataDir + "goodreads_fantasy.tsv", reader=reader)
```

code from: <https://cseweb.ucsd.edu/~jmcauley/pml/code/chap5.html>

Latent Factor Models (Surprise)

Create a model instance, train/test splits, and fit the model:

`model = SVD()`



Inbuilt functions to split into training and test fractions

```
trainset, testset = train_test_split(data, test_size=.25)
```

Fit the model and extract predictions

```
model.fit(trainset)
predictions = model.test(testset)
```

Latent Factor Models (Surprise)

Make predictions and compute MSE from the fitted model

```
predictions[0].est
```

3.6334479463688463

MSE for model predictions (test set)

```
sse = 0
for p in predictions:
    sse += (p.r_ui - p.est)**2

print(sse / len(predictions))
```

1.1883531641648757

Web Mining and Recommender Systems

Bayesian Personalized Ranking (Implicit)

Bayesian Personalized Ranking (Implicit)

A little more work to put the data in the right format. Start by reading the data in our usual formats:

```
data = list(parseData(dataDir + "goodreads_reviews_fantasy_paranormal.json.gz"))
```

```
random.shuffle(data)
```

Example from the dataset

```
data[0]
```

```
{'book_id': '13451182',  
 'date_added': 'Sun Sep 09 18:58:45 -0700 2012',  
 'date_updated': 'Sun Oct 07 15:13:32 -0700 2012',  
 'n_comments': 1,  
 'n_votes': 0,  
 'rating': 1,
```

code from:

<https://cseweb.ucsd.edu/~jmcauley/pml/code/chap5.html>

Bayesian Personalized Ranking (Implicit)

Build some utility data structures:

userIDs["julian"] → 148

```
userIDs, itemIDs = {}, {}  
  
for d in data:  
    u, i = d['user_id'], d['book_id']  
    if not u in userIDs: userIDs[u] = len(userIDs)  
    if not i in itemIDs: itemIDs[i] = len(itemIDs)  
  
nUsers, nItems = len(userIDs), len(itemIDs)
```

```
nUsers, nItems
```

```
(256088, 258212)
```

Bayesian Personalized Ranking (Implicit)

Build some sparse matrix data structures. Here we essentially build the (massive!) user-item interaction matrix describing which items users have interacted with:

```
Xiu = scipy.sparse.lil_matrix((nItems, nUsers))
for d in data:
    Xiu[itemIDs[d['book_id']],userIDs[d['user_id']]] = 1
Xui = scipy.sparse.csr_matrix(Xiu.T)
```



Bayesian Personalized Ranking model with 5 latent factors

```
model = bpr.BayesianPersonalizedRanking(factors = 5)
```

$$\sigma(\theta_u \cdot \theta_i - \theta_u \cdot \theta_j)$$

Bayesian Personalized Ranking (Implicit)

Fit the model, and get some recommendations from it:

```
model.fit(Xiu)
```

Get recommendations for a particular user (the first one) and to get items related to (similar latent factors) to a particular item

```
recommended = model.recommend(0, Xui)  
related = model.similar_items(0)
```

```
related
```

```
[(0, 1.0),  
 (42098, 0.9885355),  
 (142964, 0.9845209),  
 (150861, 0.98274595),  
 (231639, 0.9826295),  
 (182330, 0.9813926),  
 (241222, 0.98124224)]
```

Bayesian Personalized Ranking (Implicit)

Can also extract latent factors (e.g. for visualization):

```
itemFactors = model.item_factors  
userFactors = model.user_factors
```

```
itemFactors[0]
```

```
array([-0.74582803, -0.10878776,  0.32922822,  0.16516064,  0.38874012,  
       0.7460656 ], dtype=float32)
```

$\beta_i \rightarrow x_{i1} \rightarrow x_{i5}$

Web Mining and Recommender Systems

Recommender Systems in Tensorflow

Recommender Systems in Tensorflow

(will come back to Tensorflow later, but code is
in: <https://cseweb.ucsd.edu/~jmcauley/pml/code/chap5.html>)

Web Mining and Recommender Systems

More on recommender systems evaluation

Challenges in evaluating recommender systems

So far, we've mostly considered the Mean Squared Error when evaluating recommender systems; we haven't thought too hard about this since introducing linear regression

What might be some problems with this choice?

Challenges in evaluating recommender systems

What might be some problems with the MSE?

Consider e.g.



$SE = 1$

$SE = 4$

Squared
error

failed to show the
user an item they like

Which has a higher penalty? Which *should* have?

Challenges in evaluating recommender systems

What might be some problems with the MSE?

Consider e.g.

Label:					Squared Error	
Model 1:		$SE=1$		$SE=1$	$SSE=2$	
VS.						
Model 2:		$SE=\frac{1}{4}$		$SE=\frac{1}{4}$	$SSE=\frac{1}{2}$	

Model 1 gets ordering right


Which has a higher penalty? Which *should* have?

Challenges in evaluating recommender systems

What might be some problems with the MSE?
Consider e.g.

MSE assumed errors were normally distributed; what if they're more bimodal?

What *should* the correct prediction be in this case?



Challenges in evaluating recommender systems

More thoughts:

- The most popular items (or most active users) will dominate our MSE calculation; will less popular items (or users) receive "fair" consideration?
- A small change in the MSE can *drastically* change the ordering of the most relevant items; alternately a better MSE does not necessarily mean a better recommender

Ranking-based evaluation of recommender systems

Just as we saw (e.g.) precision and recall when evaluating classifiers, we can consider ranking-based metrics for evaluation of recommender systems. A few we'll look at:

- Precision and Recall @ K (again)
 - AUC (Area Under ROC Curve)
 - Mean Reciprocal Rank
- Cumulative Gain and NDCG (in textbook)
 - Beyond accuracy

Precision and Recall @ K

Much as we considered Precision and Recall (@K) when evaluating classifiers, they can also be used to evaluate ranked recommendation lists

First, rank recommended items for each user by relevance:

Lower rank =
more relevant

$$rank_u(i) < rank_u(j) \Leftrightarrow f(u, i) > f(u, j)$$

$$rank_u(i) = rank_u(j) \Leftrightarrow i = j.$$

Relevance score (e.g.
click probability)

$\sigma_u \cdot \sigma_j$

Precision and Recall @ K

Next, count how many of the (withheld/test) interactions for a user are among the top K recommendations:

$$\text{prec}_u(k) = \frac{|\{i \in I_u \mid \text{rank}_u(i) \leq k\}|}{k}$$

test set for user u

Precision and Recall @ K

Next, count how many of the (withheld/test) interactions for a user are among the top K recommendations:

$$\text{precision@}K(u) = \frac{|\{i \in I_u \mid \text{rank}_u(i) \leq K\}|}{K}$$

Test interactions



Can then be defined for *all* users (likewise for recall@K):

$$\text{precision@}K = \frac{1}{|U|} \sum_{u \in U} \text{precision@}K(u) \qquad \text{recall@}K = \frac{1}{|U|} \sum_{u \in U} \frac{|\{i \in I_u \mid \text{rank}_u(i) \leq K\}|}{|I_u|}$$

Mean Reciprocal Rank

How high is the rank of the relevant item?

- An ideal algorithm should rank it first
- An algorithm that ranks it 10th is somewhat worse
 - An algorithm that ranks it 100th is much worse
- The further down the ranking we go the less difference it makes

(assuming only a single withheld "test" item i_u for each user)

Mean Reciprocal Rank

$$MRR = \frac{1}{|Y|} \sum_u \frac{1}{\text{rank}_u(i_u)}$$

withheld item for u

1.0 = ideal algorithm; withheld item always ranked first

1/n = relevant item tends to be ranked in the n^{th} position

AUC

Does a ranker tend to give *positive* (e.g. purchased) items higher ranks than ~~*negative*~~ (e.g. not-purchased) items:


$$AUC(u) = \frac{1}{|I_u| |I \setminus I_u|} \sum_{i \in I_u} \sum_{\substack{j \notin I_u \\ j \in I \setminus I_u}} \delta(\text{rank}_u(i) < \text{rank}_u(j))$$

non
- interaction

$$\delta(\sigma_u \cdot \sigma_i > \sigma_u \cdot \sigma_j)$$

AUC

$$\sum_i \sum_j \delta(\text{rank}_u(i) < \text{rank}_u(j))$$

$$\sim \sum_i \sum_j \sigma(\underbrace{\delta_u \cdot x_i - \delta_u \cdot x_j}_{\text{---}}) > 0.5$$


$$\sim \max_{\delta} \sum_i \sum_j \log \sigma(\delta_u \cdot x_i - \delta_u \cdot x_j)$$

AUC

The AUC:

- Counts the fraction of times the algorithm gives a *higher score to a positive than to a negative interaction*
 - (1.0 = always correct; 0.5 = random)
 - Across all users:

$$\text{AUC} = \frac{1}{|U|} \text{AUC}(u)$$

AUC

Why the AUC?

- ① Reasonable ranking function to optimize
- ② Easy to optimize

Why the AUC?

- Doesn't force negative items to be rated as "negative" – just less positive than positive – this is desirable in implicit feedback settings
- Rewards the algorithm for *ordering things correctly*, which may be more important to users than (e.g.) predicting ratings correctly
- Is easy (compared to some other metrics) to optimize

Web Mining and Recommender Systems

"User-free" models of recommendation

User-free recommenders

So far we've studied (arguably) the simplest approaches for a variety of tasks:

1. Similarity-based models for **binary** data
2. Similarity-based models for **real-valued** data
3. Dimensionality reduction (latent factors) for **real-valued** data
4. Dimensionality reduction for **binary** data

Next we'll discuss a few alternate approaches to similar problems

User-free recommenders

Main goal in this section is to *avoid having an explicit model of a user* γ_u

Why?

- It's costly (memory) $|K| \times |U|$
- It's hard to update
 - new users
 - updating γ_u for a single click

User-free recommenders

Main goal in this section is to *avoid having an explicit model of a user γ_u*

Why?

- Previous models had K parameters per user – very expensive in settings with many users!
- Poor performance for users with few interactions
- Requires continuous retraining as users continue to interact

User-free recommenders

Can we design algorithms that take a *set of items* as input, and generate recommended items as output? $f(\text{history}(u), i)$

As users provide more interactions, we just change the input to the algorithm – no need to retrain!

1. Sparse Linear Methods (SLIM)
2. Factored Item Similarity Models (FISM)


1. Sparse Linear Methods (SLIM) (Ning and Karypis, 2011)

Adapts ideas from regression; model interactions as

$$f(u, i) = R_u \cdot W_i$$

Vector of
interactions
for user u

(linear!)
parameter
vector



$$f(u, i) = \frac{1}{\sum_{j \in I_u} S_{u,j}} \sum_{j \in I_u} r_{u,j} \times \cancel{S_{u,j}(i,j)} \cdot W_{ij}$$

1. Sparse Linear Methods (SLIM) (Ning and Karypis, 2011)

Can be expanded as:

$$f(u, i) = \sum_{j \in I_u} R_{uj} \cdot w_{ij}$$
$$W = \begin{bmatrix} w_{0,0} & w_{0,1} & \dots & w_{0,I-1} \\ \vdots & & & \\ w_{I-1,0} & & & \end{bmatrix}$$

Challenge: W is a (dense!) $|I| \times |I|$ matrix

1. Sparse Linear Methods (SLIM) (Ning and Karypis, 2011)

Solution: assume W is *sparse*, which is achieved through a regularizer:

$$\begin{aligned} \arg \min_W & \|R - RW^T\|_2^2 + \lambda \Omega_2(W) + \lambda' \Omega_1(W) \\ \text{s.t.} \quad & W_{i,j} \geq 0; \quad W_{i,i} = 0. \end{aligned}$$

1. Sparse Linear Methods (SLIM) (Ning and Karypis, 2011)

Sparse linear methods:

- Rapid inference time (compared to traditional methods, not compared to latent factor models)
- Better long-tail performance, i.e., works well for rarely-occurring items

2. Factored Item Similarity Models (FISM) (Kabbur et al. 2013)

Idea: a user is just the average over items they consume

Replace the user representation with an average of item representations

$$f(u, i) = \alpha + \beta_u + \beta_i + \gamma_u \cdot \gamma_i$$

2. Factored Item Similarity Models (FISM) (Kabbur et al. 2013)

Replace

$$f(u, i) = \alpha + \beta_u + \beta_i + \gamma_u \cdot \gamma_i$$

not quite "user-free"

with

$$f(u, i) = \alpha + \beta_u + \beta_i +$$

$$\sum_{j \in I_u} \gamma_j' \cdot \gamma_i$$

2. Factored Item Similarity Models (FISM) (Kabbur et al. 2013)

Replace

$$f(u, i) = \alpha + \beta_u + \beta_i + \gamma_u \cdot \gamma_i$$

with

$$f(u, i) = \alpha + \beta_u + \beta_i + \frac{1}{|I_u \setminus \{i\}|} \sum_{j \in I_u \setminus \{i\}} \gamma'_j \cdot \gamma_i$$

$|U| \times K + \underbrace{|I| \times K}_{\substack{\text{params}}}$
 $2|I|K$
params

2. Factored Item Similarity Models (FISM) (Kabbur et al. 2013)

Note that we have *two* item representations (instead of an item and a user representation)

γ'_j and γ_i

These are essentially a "query" and a "target" representation

References

Further reading:

- Sparse Linear Methods (SLIM) (Ning and Karypis, 2011)
- Factored Item Similarity Models (FISM) (Kabbur et al. 2013)

Web Mining and Recommender Systems

Deep learning for recommendation

Deep learning for recommendation

Won't spend a tonne of time teaching deep learning, but obviously it's made some headway into recommendation (just like everywhere else...)

Here will just give a basic sense of some of the main ideas

See textbook for details!



Why *not* deep learning for recommendation?



Why *should* we need deep learning to improve recommender systems?

$$f(u, i) = \alpha + \beta_u + \beta_i + \gamma_u \cdot \gamma_i$$

Deep Learning is normally used to uncover non-linear relationships/transforms among features, but this model doesn't *have* any features!

If latent factors can uncover any properties about users/items, what can a "deep" model learn?





Why *not* deep learning for recommendation?


Idea: there's nothing sacred about the inner product in this function, and other choices might be better in some contexts

$$f(u, i) = \alpha + \beta_u + \beta_i + \gamma_u \cdot \gamma_i$$

Maybe we can automatically learn what types of relationship would be most effective


1. Neural Collaborative Filtering (He et al. 2017)

Idea: use a multilayer perceptron to learn the ideal relationship between γ_u and γ_i



1. Neural Collaborative Filtering (He et al. 2017)

Idea: use a multilayer perceptron to learn the ideal relationship between γ_u and γ_i




1. Neural Collaborative Filtering (He et al. 2017)

Idea: use a multilayer perceptron to learn the ideal relationship between γ_u and γ_i

- Arguably, this will help us to learn more complex interactions between users and items
 - Some counterargument (e.g. Dacrema et al. 2019): it's actually hard for an MLP to learn an inner product function!


2. AutoRec (Sedhain et al. 2015)

Idea: *Autoencoders* are a technique to learn low-dimensional latent representations of feature vectors; can they be used for recommendation?



2. AutoRec (Sedhain et al. 2015)

Idea: *Autoencoders* are a technique to learn low-dimensional latent representations of feature vectors; can they be used for recommendation?



2. AutoRec (Sedhain et al. 2015)

Autoencoder-based recommendation:

- Input is a vector of items a user has consumed (or a set of users who have consumed an item)
 - Model is trained to encode these vectors
- At test time, find un-consumed items that have the highest score according to the decoder
 - **Note:** this is a user-free model! The input/output is just a vector of items!

3. Convolutional and Recurrent Models

Plenty of other approaches based on Convolutional and Recurrent Neural Networks:

- CNNs often used as a way to incorporate rich *content* into recommender systems (e.g. images)
- RNNs (and Transformers, etc.) are often used as a way to incorporate sequential dynamics

We'll come back to these a little (but not much) later

Are deep-learning models "worth it"?

Some doubts as to whether deep learning-based recommenders are really "worth it":

- Some evidence that simpler models will work just as well if carefully tuned (Dacrema et al. 2019)
- Potentially adding a lot of parameters / training complexity for modest performance improvements
 - Also challenges re. efficient retrieval etc.

References

Further reading:

- He et al. (2017): Neural Collaborative Filtering
- He and Chua (2017): Neural Factorization Machines
- Cheng et al. (2016): Wide and Deep Learning for Recommendation
- Guo et al. (2017): Deep Factorization Machines
 - Sedhain et al. (2015): AutoRec

Web Mining and Recommender Systems

What's still coming up?

Extensions of latent-factor models

So far we have a model that looks like:

$$f(u, i) = \alpha + \beta_u + \beta_i + \gamma_u \cdot \gamma_i$$

How might we extend this to:

- Incorporate features about users and items
 - Handle implicit feedback
 - Change over time

See **Yehuda Koren** (+Bell & Volinsky)'s magazine article:
"Matrix Factorization Techniques for Recommender Systems"
IEEE Computer, 2009

Extensions of latent-factor models

1) Features about users and/or items

(simplest case) Suppose we have **binary attributes** to describe users or items

$$A(u) = [1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 1]$$

attribute vector for user u

e.g. is female is male is between 18-24yo

The diagram illustrates the attribute vector $A(u)$ for a user u . The vector is a 1x12 array of binary values: $[1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 1]$. Three green arrows point from descriptive text below to specific elements in the vector: the first arrow points from 'e.g. is female' to the second element (0), the second arrow points from 'is male' to the third element (1), and the third arrow points from 'is between 18-24yo' to the tenth element (1). A fourth green arrow points from 'attribute vector for user u ' to the first element (1).

Extensions of latent-factor models

1) Features about users and/or items

(simplest case) Suppose we have **binary attributes** to describe users or items

- Associate a **parameter vector** with each attribute
- Each vector encodes how much a particular feature "offsets" the given latent dimensions

$$A(u) = [1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 1]$$

attribute vector for user u

e.g. $y_0 = [-0.2, 0.3, 0.1, -0.4, 0.8]$
~ "how does being male impact γ_u "

Extensions of latent-factor models

1) Features about users and/or items

(simplest case) Suppose we have **binary attributes** to describe users or items

- Associate a **parameter vector** with each attribute
- Each vector encodes how much a particular feature "offsets" the given latent dimensions
 - Model looks like:

$$f(u, i) = \alpha + \beta_u + \beta_i + \left(\gamma_u + \sum_{a \in A(u)} \rho_a \right) \cdot \gamma_i$$

- Fit as usual:

$$\arg \min_{\alpha, \beta, \gamma, \rho} \sum_{u, i \in \text{train}} \underbrace{(f(u, i) - r_{u, i})^2}_{\text{error}} + \underbrace{\lambda \Omega(\beta, \gamma)}_{\text{regularizer}}$$

Extensions of latent-factor models

2) Implicit feedback

Perhaps many users will never actually rate things, but may still interact with the system, e.g. through the movies they view, or the products they purchase (but never rate)

- Adopt a similar approach – introduce a binary vector describing a user's actions

$$N(u) = [1, 0, 0, 0, 1, 0, \dots, 0, 1]$$

implicit feedback vector for user u

e.g. $y_u = [-0.1, 0.2, 0.3, -0.1, 0.5]$

Clicked on "Love Actually" but didn't watch

Extensions of latent-factor models

2) Implicit feedback

Perhaps many users will never actually rate things, but may still interact with the system, e.g. through the movies they view, or the products they purchase (but never rate)

- Adopt a similar approach – introduce a binary vector describing a user's actions
 - Model looks like:

$$f(u, i) = \alpha + \beta_u + \beta_i + \left(\gamma_u + \frac{1}{\|N(u)\|} \sum_{a \in N(u)} \rho_a \right) \cdot \gamma_i$$

normalize by the number of actions the user performed

Extensions of latent-factor models


3) Change over time

There are a number of reasons why rating data might be subject to temporal effects...

- seasonal, weekly, daily
- sequential dynamics
- age of items


Extensions of latent-factor models

3) Change over time




Extensions of latent-factor models

3) Change over time



Extensions of latent-factor models

3) Change over time




A few temporal effects from beer reviews

Extensions of latent-factor models

3) Change over time

There are a number of reasons why rating data might be subject to temporal effects...



Extensions of latent-factor models

3) Change over time

Each definition of temporal evolution demands a slightly different model assumption (we'll see some in more detail later tonight!) but the basic idea is the following:

1) Start with our original model:

$$f(u, i) = \alpha + \beta_u + \beta_i + \gamma_u \cdot \gamma_i$$

2) And define some of the parameters as a function of time:

$$f(u, i, t) = \alpha + \beta_u(t) + \beta_i(t) + \gamma_u(t) \cdot \gamma_i$$

3) Add a regularizer to constrain the time-varying terms:


$$\arg \min_{\alpha, \beta, \gamma} \sum_{u, i, t \in \text{train}} (f(u, i, t) - r_{u, i, t})^2 + \lambda_1 \Omega(\beta, \gamma) + \underbrace{\lambda_2 \|\gamma(t) - \gamma(t + \delta)\|}_{\text{parameters should change smoothly}}$$

parameters should change smoothly

Moral(s) of the story

How much do these extensions help?

Moral: increasing complexity helps a bit, but changing the model can help **a lot**



Moral(s) of the story

So what actually happened with Netflix?

- The AT&T team “BellKor”, consisting of Yehuda Koren, Robert Bell, and Chris Volinsky were early leaders. Their main insight was how to effectively incorporate temporal dynamics into recommendation on Netflix.
- Before long, it was clear that no one team would build the winning solution, and Frankenstein efforts started to merge. Two frontrunners emerged, “BellKor’s Pragmatic Chaos”, and “The Ensemble”.
- The BellKor team was the first to achieve a 10% improvement in RMSE, putting the competition in “last call” mode. The winner would be decided after 30 days.
- After 30 days, performance was evaluated on the hidden part of the test set.
- Both of the frontrunning teams had **the same** RMSE (up to some precision) but BellKor’s team submitted their solution 20 minutes earlier and won \$1,000,000

For a less rough summary, see the Wikipedia page about the Netflix prize, and the nytimes article about the competition: <http://goo.gl/WNpy7o>

Moral(s) of the story

Afterword

- Netflix had a class-action lawsuit filed against them after somebody de-anonymized the competition data
- \$1,000,000 seems to be **incredibly cheap** for a company the size of Netflix in terms of the amount of research that was devoted to the task, and the potential benefit to Netflix of having their recommendation algorithm improved by 10%
- Other similar competitions have emerged, such as the Heritage Health Prize (\$3,000,000 to predict the length of future hospital visits)
- But... the winning solution never made it into production at Netflix – it's a monolithic algorithm that is very expensive to update as new data comes in*

*source: a friend of mine told me and I have no actual evidence of this claim

Moral(s) of the story

Finally...

Q: Is the RMSE really the right approach? Will improving rating prediction by 10% actually improve the user experience by a significant amount?

A: Not clear. Even a solution that only changes the RMSE slightly could drastically change which items are top-ranked and ultimately suggested to the user.

Q: But... are the following recommendations actually any good?

A1: Yes, these are my favorite movies!

or **A2:** No! There's no **diversity**, so how will I discover **new** content?




5.0 stars



5.0 stars



5.0 stars



5.0 stars



4.9 stars



4.9 stars



4.8 stars



4.8 stars

predicted rating

Summary

Various extensions of latent factor models:

- Incorporating features
e.g. for cold-start recommendation
 - Implicit feedback
e.g. when ratings aren't available, but other actions are
- Incorporating temporal information into latent factor models
seasonal effects, short-term "bursts", long-term trends, etc.
 - Missing-not-at-random
incorporating priors about items that were not bought or rated
 - The Netflix prize

References

Further reading:

Yehuda Koren's, Robert Bell, and Chris Volinsky's IEEE computer article:

<http://www2.research.att.com/~volinsky/papers/ieeecomputer.pdf>

Paper about the "Missing-at-Random" assumption, and how to address it:

<http://www.cs.toronto.edu/~marlin/research/papers/cfmar-uai2007.pdf>

Collaborative filtering with temporal dynamics:

<http://research.yahoo.com/files/kdd-fp074-koren.pdf>

Recommender systems and sales diversity:

http://papers.ssrn.com/sol3/papers.cfm?abstract_id=955984